

- (1) Differentiation of Univariate Functions
- (2) Partial Differentiation and Gradients
- (3) Gradients of Vector-Valued Functions
- (4) Gradients of Matrices
- (5) Useful Identities for Computing Gradients
- (6) Backpropagation and Automatic Differentiation
- (7) Higher-Order Derivatives
- (8) Linearization and Multivariate Taylor Series

- ▶ Machine learning is about solving an optimization problem whose variables are the parameters of a given model.
- ▶ Solving optimization problems requires gradient information.
- ▶ Central to this chapter is the concept of the function, which we often write

$$f : \mathbb{R}^D \mapsto \mathbb{R}$$

$$\mathbf{x} \mapsto f(\mathbf{x})$$

- (1) Differentiation of Univariate Functions
- (2) Partial Differentiation and Gradients
- (3) Gradients of Vector-Valued Functions
- (4) Gradients of Matrices
- (5) Useful Identities for Computing Gradients
- (6) Backpropagation and Automatic Differentiation
- (7) Higher-Order Derivatives
- (8) Linearization and Multivariate Taylor Series

# Why Should I Care About Derivatives?

- ▶ **Training AI models:** Every time a neural network learns, it computes derivatives to figure out how to improve
- ▶ **Gradient descent:** The derivative tells the model which direction to “walk” to reduce its errors
- ▶ **Physics simulations:** Derivatives describe velocity, acceleration, and change in games and simulations
- ▶ **Economics:** Marginal cost, marginal revenue — all derivatives!

**In one sentence:** Derivatives tell you how things change — and AI is all about learning from change.

# Difference Quotient and Derivative

- ▶ **Difference Quotient.** The average slope of  $f$  between  $x$  and  $x + \partial x$

$$\frac{\partial y}{\partial x} \stackrel{\text{def}}{=} \frac{f(x + \partial x) - f(x)}{\partial x}$$

- ▶ **Derivative.** Pointing in the direction of steepest ascent of  $f$ .

$$\frac{df}{dx} \stackrel{\text{def}}{=} \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

- ▶ Unless confusion arises, we often use  $f' = \frac{df}{dx}$ .

# Taylor Series

- ▶ Representation of a function as an infinite sum of terms, using derivatives of  $f$  evaluated at  $x_0$ .
- ▶ **Taylor polynomial.** The Taylor polynomial of degree  $n$  of  $f : \mathbb{R} \mapsto \mathbb{R}$  at  $x_0$  is:

$$T_n(x) \stackrel{\text{def}}{=} \sum_{k=0}^n \frac{f^{(k)}(x_0)}{k!} (x - x_0)^k, \text{ where } f^{(k)}(x_0) \text{ is the } k\text{th derivative of } f \text{ at } x_0.$$

- ▶ **Taylor Series.** For a smooth function  $f \in \{C\}^\infty$ , the Taylor series of  $f$  at  $x_0$  is:

$$T_\infty(x) \stackrel{\text{def}}{=} \sum_{k=0}^{\infty} \frac{f^{(k)}(x_0)}{k!} (x - x_0)^k.$$

- ▶ If  $f(x) = T_\infty(x)$ ,  $f$  is called **analytic**.

# Differentiation Rules

- ▶ Product rule.  $(f(x)g(x))' = f'(x)g(x) + f(x)g'(x)$
- ▶ Quotient rule.  $\left(\frac{f(x)}{g(x)}\right)' = \frac{f'(x)g(x) - f(x)g'(x)}{(g(x))^2}$
- ▶ Sum rule.  $(f(x) + g(x))' = f'(x) + g'(x)$
- ▶ Chain rule.  $(g(f(x)))' = g'(f(x))f'(x)$

## What we just learned:

- ▶ **Derivative** = slope of a function at a point (rate of change)
- ▶ **Taylor series** = approximate any smooth function using its derivatives
- ▶ Key rules: **product, quotient, sum, chain** — these are your tools for computing derivatives
- ▶ **Next up:** What happens when functions have *multiple* inputs?

- (1) Differentiation of Univariate Functions
- (2) Partial Differentiation and Gradients
- (3) Gradients of Vector-Valued Functions
- (4) Gradients of Matrices
- (5) Useful Identities for Computing Gradients
- (6) Backpropagation and Automatic Differentiation
- (7) Higher-Order Derivatives
- (8) Linearization and Multivariate Taylor Series

# Why Should I Care About Gradients?

- ▶ **Multi-dimensional optimization:** Real models have thousands to billions of parameters — you need **partial derivatives** for each one
- ▶ **The gradient** points in the direction of steepest ascent (negate it → steepest descent = learning!)
- ▶ **Image recognition:** The gradient tells you which pixels matter most for the model's prediction
- ▶ **ChatGPT & LLMs:** Every word prediction involves gradient-based updates across millions of parameters

**In one sentence:** The gradient is the compass that guides AI models toward better predictions.

# Gradient

- ▶ Now,  $f : \mathbb{R}^n \mapsto \mathbb{R}$ .
- ▶ Gradient of  $f$  w.r.t.  $\mathbf{x}$ , denoted  $\nabla_{\mathbf{x}} f$ : Vary one variable at a time, keeping the others constant.

**Partial Derivative.** For  $f : \mathbb{R}^n \mapsto \mathbb{R}$ ,

$$\frac{\partial f}{\partial x_1} = \lim_{h \rightarrow 0} \frac{f(x_1 + h, x_2, \dots, x_n) - f(\mathbf{x})}{h}$$

⋮

$$\frac{\partial f}{\partial x_n} = \lim_{h \rightarrow 0} \frac{f(x_1, x_2, \dots, x_n + h) - f(\mathbf{x})}{h}$$

**Gradient.** Get the partial derivatives and collect them in the row vector.

$$\nabla_{\mathbf{x}} f = \frac{df}{d\mathbf{x}} = \begin{pmatrix} \frac{\partial f(\mathbf{x})}{\partial x_1} & \dots & \frac{\partial f(\mathbf{x})}{\partial x_n} \end{pmatrix} \in \mathbb{R}^{1 \times n}$$

## Example

► **Example.**  $f(x, y) = (x + 2y^3)^2$

$$\frac{\partial f(x, y)}{\partial x} = 2(x + 2y^3) \frac{\partial x + 2y^3}{\partial x} = 2(x + 2y^3)$$

$$\frac{\partial f(x, y)}{\partial y} = 2(x + 2y^3) \frac{\partial x + 2y^3}{\partial y} = 12(x + 2y^3)y^2$$

► **Example.**  $f(x_1, x_2) = x_1^2 x_2 + x_1 x_2^3$

$$\nabla_{(x_1, x_2)} f = \frac{df}{dx} = \begin{pmatrix} \frac{\partial f(x_1, x_2)}{\partial x_1} & \frac{\partial f(x_1, x_2)}{\partial x_2} \end{pmatrix} = \begin{pmatrix} 2x_1 x_2 + x_2^3 & x_1^2 + 3x_1 x_2^2 \end{pmatrix}$$

# Rules for Partial Differentiation

- ▶ Product rule

$$\frac{\partial}{\partial \mathbf{x}}(f(\mathbf{x})g(\mathbf{x})) = \frac{\partial f}{\partial \mathbf{x}}g(\mathbf{x}) + f(\mathbf{x})\frac{\partial g}{\partial \mathbf{x}}$$

- ▶ Sum rule

$$\frac{\partial}{\partial \mathbf{x}}(f(\mathbf{x}) + g(\mathbf{x})) = \frac{\partial f}{\partial \mathbf{x}} + \frac{\partial g}{\partial \mathbf{x}}$$

- ▶ Chain rule

$$\frac{\partial}{\partial \mathbf{x}}g(f(\mathbf{x})) = \frac{\partial g}{\partial f} \frac{\partial f}{\partial \mathbf{x}}$$

## More about Chain Rule

- $f : \mathbb{R}^2 \mapsto \mathbb{R}$  of two variables  $x_1$  and  $x_2$ .  $x_1(t)$  and  $x_2(t)$  are functions of  $t$ .

$$\frac{df}{dt} = \begin{pmatrix} \frac{\partial f}{\partial x_1} & \frac{\partial f}{\partial x_2} \end{pmatrix} \begin{pmatrix} \frac{\partial x_1(t)}{\partial t} \\ \frac{\partial x_2(t)}{\partial t} \end{pmatrix} = \frac{\partial f}{\partial x_1} \frac{\partial x_1}{\partial t} + \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial t}$$

- **Example.**  $f(x_1, x_2) = x_1^2 + 2x_2$ , where  $x_1(t) = \sin(t)$ ,  $x_2(t) = \cos(t)$

$$\frac{df}{dt} = \frac{\partial f}{\partial x_1} \frac{\partial x_1}{\partial t} + \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial t} = 2 \sin(t) \cos(t) - 2 \sin(t) = 2 \sin(t)(\cos(t) - 1)$$

- $f : \mathbb{R}^2 \mapsto \mathbb{R}$  of two variables  $x_1$  and  $x_2$ .  $x_1(s, t)$  and  $x_2(s, t)$  are functions of  $s, t$ .

$$\begin{aligned}\frac{\partial f}{\partial s} &= \frac{\partial f}{\partial x_1} \frac{\partial x_1}{\partial s} + \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial s} \\ \frac{\partial f}{\partial t} &= \frac{\partial f}{\partial x_1} \frac{\partial x_1}{\partial t} + \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial t}\end{aligned}$$

$$\frac{df}{d(s, t)} = \frac{\partial f}{\partial \mathbf{x}} \frac{\partial \mathbf{x}}{\partial (s, t)} = \begin{pmatrix} \frac{\partial f}{\partial x_1} & \frac{\partial f}{\partial x_2} \end{pmatrix} \begin{pmatrix} \frac{\partial x_1}{\partial s} & \frac{\partial x_1}{\partial t} \\ \frac{\partial x_2}{\partial s} & \frac{\partial x_2}{\partial t} \end{pmatrix}$$

# Quick Recap: Partial Derivatives & Gradients

## What we just learned:

- ▶ **Partial derivative:** change one variable, keep the rest fixed
- ▶ **Gradient:** collect all partial derivatives into a row vector
- ▶ **Chain rule** extends naturally to multiple variables
- ▶ **Next up:** What if the output is also a vector? → Jacobian matrices

- (1) Differentiation of Univariate Functions
- (2) Partial Differentiation and Gradients
- (3) **Gradients of Vector-Valued Functions**
- (4) Gradients of Matrices
- (5) Useful Identities for Computing Gradients
- (6) Backpropagation and Automatic Differentiation
- (7) Higher-Order Derivatives
- (8) Linearization and Multivariate Taylor Series

$$\mathbf{f} : \mathbb{R}^n \mapsto \mathbb{R}^m$$

- ▶ For a function  $\mathbf{f} : \mathbb{R}^n \mapsto \mathbb{R}^m$  and vector  $\mathbf{x} = (x_1 \ \dots \ x_n)^\top \in \mathbb{R}^n$ , the vector-valued function is:

$$\mathbf{f}(\mathbf{x}) = \begin{pmatrix} f_1(\mathbf{x}) \\ \vdots \\ f_m(\mathbf{x}) \end{pmatrix}$$

- ▶ Partial derivative w.r.t.  $x_i$  is a column vector:  $\frac{\partial \mathbf{f}}{\partial x_i} = \begin{pmatrix} \frac{\partial f_1}{\partial x_i} \\ \vdots \\ \frac{\partial f_m}{\partial x_i} \end{pmatrix}$
- ▶ Gradient (or Jacobian):  $\frac{d\mathbf{f}(\mathbf{x})}{d\mathbf{x}} = \left( \frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_1} \ \dots \ \frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_n} \right)$

# Jacobian

$$\begin{aligned}\mathbf{J} = \nabla_{\mathbf{x}} \mathbf{f} &= \frac{d\mathbf{f}(\mathbf{x})}{d\mathbf{x}} = \begin{pmatrix} \frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_1} & \dots & \frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_n} \end{pmatrix} \\ &= \begin{pmatrix} \frac{\partial f_1(\mathbf{x})}{\partial x_1} & \dots & \frac{\partial f_1(\mathbf{x})}{\partial x_n} \\ \vdots & & \vdots \\ \frac{\partial f_m(\mathbf{x})}{\partial x_1} & \dots & \frac{\partial f_m(\mathbf{x})}{\partial x_n} \end{pmatrix}\end{aligned}$$

- For a  $\mathbb{R}^n \rightarrow \mathbb{R}^m$  function, its Jacobian is an  $m \times n$  matrix.

## Example: Gradient of Vector-Valued Function

- ▶  $\mathbf{f}(\mathbf{x}) = \mathbf{Ax}$ ,  $\mathbf{f} : \mathbb{R}^n \mapsto \mathbb{R}^m$ ,  $\mathbf{A} \in \mathbb{R}^{m \times n}$ ,  $\mathbf{x} \in \mathbb{R}^n$
- ▶ Partial derivatives:  $f_i(\mathbf{x}) = \sum_{j=1}^n A_{ij}x_j \implies \frac{\partial f_i}{\partial x_j} = A_{ij}$
- ▶ Gradient

$$\frac{d\mathbf{f}}{d\mathbf{x}} = \begin{pmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{pmatrix} = \begin{pmatrix} A_{11} & \cdots & A_{1n} \\ \vdots & & \vdots \\ A_{m1} & \cdots & A_{mn} \end{pmatrix} = \mathbf{A}$$

## Example: Chain Rule

- $h : \mathbb{R} \mapsto \mathbb{R}$ ,  $h(t) = (f \circ g)(t)$  with

$$f : \mathbb{R}^2 \mapsto \mathbb{R}, \quad f(\mathbf{x}) = \exp(x_1 x_2^2), \quad g : \mathbb{R} \mapsto \mathbb{R}^2, \quad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = g(t) = \begin{pmatrix} t \cos(t) \\ t \sin(t) \end{pmatrix}$$

- (Note)  $\frac{\partial f}{\partial \mathbf{x}} \in \mathbb{R}^{1 \times 2}$  and  $\frac{\partial g}{\partial t} \in \mathbb{R}^{2 \times 1}$
- Using the chain rule,

$$\begin{aligned}\frac{dh}{dt} &= \frac{\partial f}{\partial \mathbf{x}} \frac{\partial \mathbf{x}}{\partial t} = \left( \frac{\partial f}{\partial x_1} \quad \frac{\partial f}{\partial x_2} \right) \begin{pmatrix} \frac{\partial x_1}{\partial t} \\ \frac{\partial x_2}{\partial t} \end{pmatrix} \\ &= (\exp(x_1 x_2^2) x_2^2 \quad 2 \exp(x_1 x_2^2) x_1 x_2) \begin{pmatrix} \cos(t) - t \sin(t) \\ \sin(t) + t \cos(t) \end{pmatrix}\end{aligned}$$

## Example: Least-Square Loss (1)

- ▶ A linear model:  $\mathbf{y} = \Phi\theta$
- ▶  $\theta \in \mathbb{R}^D$ : parameter vector
- ▶  $\Phi \in \mathbb{R}^{N \times D}$ : input features
- ▶  $\mathbf{y} \in \mathbb{R}^N$ : observations
- ▶ Goal: Find a good parameter vector that provides the best-fit, formulated by minimizing the following loss  $L : \mathbb{R}^D \mapsto \mathbb{R}$  over the parameter vector  $\theta$ .

$$L(\mathbf{e}) \stackrel{\text{def}}{=} \|\mathbf{e}\|^2, \quad \text{where } \mathbf{e}(\theta) = \mathbf{y} - \Phi\theta$$

## Example: Least-Square Loss (2)

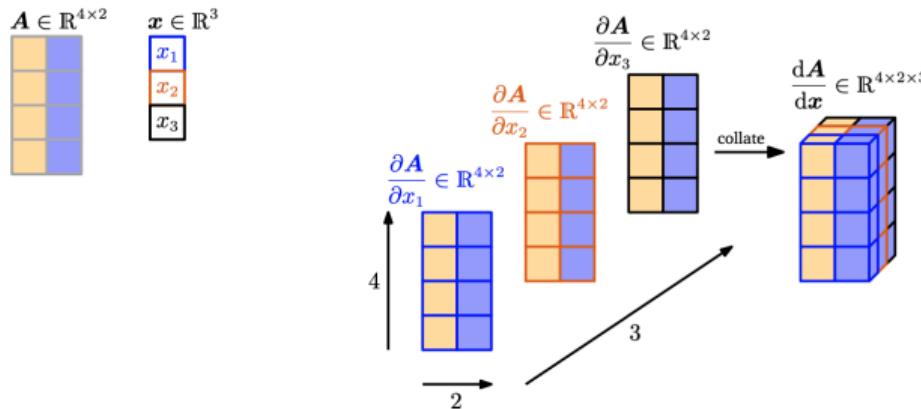
- ▶  $\frac{\partial L}{\partial \theta} = \frac{\partial L}{\partial \mathbf{e}} \frac{\partial \mathbf{e}}{\partial \theta}$
- ▶ Note.  $\frac{\partial L}{\partial \theta} \in \mathbb{R}^{1 \times D}$ ,  $\frac{\partial L}{\partial \mathbf{e}} \in \mathbb{R}^{1 \times N}$ ,  $\frac{\partial \mathbf{e}}{\partial \theta} \in \mathbb{R}^{N \times D}$
- ▶ Using that  $\|\mathbf{e}\|^2 = \mathbf{e}^\top \mathbf{e}$ ,  $\frac{\partial L}{\partial \mathbf{e}} = 2\mathbf{e}^\top \in \mathbb{R}^{1 \times N}$  and  $\frac{\partial \mathbf{e}}{\partial \theta} = -\Phi \in \mathbb{R}^{N \times D}$

Finally, we get:  $\frac{\partial L}{\partial \theta} = 2\mathbf{e}^\top (-\Phi) = -\underbrace{2(\mathbf{y}^\top - \theta^\top \Phi^\top)}_{1 \times N} \underbrace{\Phi}_{N \times D}$

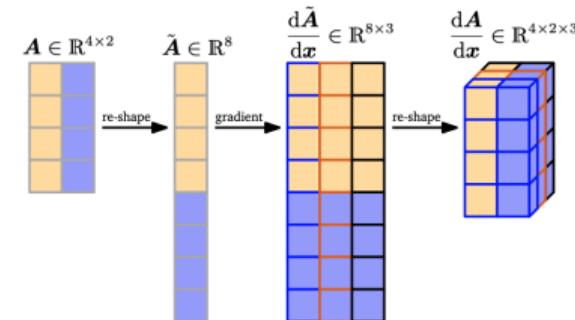
- (1) Differentiation of Univariate Functions
- (2) Partial Differentiation and Gradients
- (3) Gradients of Vector-Valued Functions
- (4) **Gradients of Matrices**
- (5) Useful Identities for Computing Gradients
- (6) Backpropagation and Automatic Differentiation
- (7) Higher-Order Derivatives
- (8) Linearization and Multivariate Taylor Series

# Gradients of Matrices

- ▶ Gradient of matrix  $\mathbf{A} \in \mathbb{R}^{m \times n}$  w.r.t. matrix  $\mathbf{B} \in \mathbb{R}^{p \times q}$
- ▶ Jacobian: A four-dimensional tensor<sup>1</sup>  $\mathbf{J} = \frac{d\mathbf{A}}{d\mathbf{B}} \in \mathbb{R}^{(m \times n) \times (p \times q)}$



(a) Approach 1: We compute the partial derivative  $\frac{\partial \mathbf{A}}{\partial x_1}, \frac{\partial \mathbf{A}}{\partial x_2}, \frac{\partial \mathbf{A}}{\partial x_3}$ , each of which is a  $4 \times 2$  matrix, and collate them in a  $4 \times 2 \times 3$  tensor.



(b) Approach 2: We re-shape (flatten)  $\mathbf{A} \in \mathbb{R}^{4 \times 2}$  into a vector  $\tilde{\mathbf{A}} \in \mathbb{R}^8$ . Then, we compute the gradient  $\frac{d\tilde{\mathbf{A}}}{d\mathbf{x}} \in \mathbb{R}^{8 \times 3}$ . We obtain the gradient tensor by re-shaping this gradient as illustrated above.

<sup>1</sup>A multidimensional array

## Example: Gradient of Vectors for Matrices (1)

- $f(\mathbf{x}) = \mathbf{Ax}$ ,  $\mathbf{f} \in \mathbb{R}^m$ ,  $\mathbf{A} \in \mathbb{R}^{m \times n}$ ,  $\mathbf{x} \in \mathbb{R}^n$ . What is  $\frac{d\mathbf{f}}{d\mathbf{A}}$ ?
- Dimension: If we consider  $\mathbf{f} : \mathbb{R}^{m \times n} \mapsto \mathbb{R}^m$ ,  $\frac{d\mathbf{f}}{d\mathbf{A}} \in \mathbb{R}^{m \times (m \times n)}$

► Partial derivatives:  $\frac{\partial f_i}{\partial \mathbf{A}} \in \mathbb{R}^{1 \times (m \times n)}$ ,  $\frac{d\mathbf{f}}{d\mathbf{A}} = \begin{pmatrix} \frac{\partial f_1}{\partial \mathbf{A}} \\ \vdots \\ \frac{\partial f_m}{\partial \mathbf{A}} \end{pmatrix}$

$$f_i = \sum_{j=1}^n A_{ij}x_j, \quad i = 1, \dots, m \implies \frac{\partial f_i}{\partial A_{iq}} = x_q,$$

$$\frac{\partial f_i}{\partial A_{i \cdot}} = \mathbf{x}^\top \in \mathbb{R}^{1 \times 1 \times n} \text{ (for } i\text{th row vector)}$$

$$\frac{\partial f_i}{\partial A_{k \neq i \cdot}} = \mathbf{0}^\top \in \mathbb{R}^{1 \times 1 \times n} \text{ (for } k\text{th row vector, } k \neq i)$$

$$\frac{\partial f_i}{\partial \mathbf{A}} = \begin{pmatrix} \mathbf{0}^\top \\ \vdots \\ \mathbf{0}^\top \\ \mathbf{x}^\top \\ \mathbf{0}^\top \\ \vdots \\ \mathbf{0}^\top \end{pmatrix} \in \mathbb{R}^{1 \times (m \times n)}$$

## Example: Gradient of Matrices for Matrices (2)

- $\mathbf{R} \in \mathbb{R}^{m \times n}$  and  $\mathbf{f} : \mathbb{R}^{m \times n} \mapsto \mathbb{R}^{n \times n}$  with  $\mathbf{f}(\mathbf{R}) = \mathbf{K} \stackrel{\text{def}}{=} \mathbf{R}^\top \mathbf{R} \in \mathbb{R}^{n \times n}$ . What is  $\frac{d\mathbf{K}}{d\mathbf{R}} \in \mathbb{R}^{(n \times n) \times (m \times n)}$ ?
- $\frac{dK_{pq}}{d\mathbf{R}} \in \mathbb{R}^{1 \times m \times n}$ . Let  $\mathbf{r}_i$  be the  $i$ th column of  $\mathbf{R}$ . Then  $K_{pq} = \mathbf{r}_p^\top \mathbf{r}_q = \sum_{k=1}^m R_{kp} R_{kq}$ .
- Partial derivative  $\frac{\partial K_{pq}}{\partial R_{ij}}$

$$\frac{\partial K_{pq}}{\partial R_{ij}} = \sum_{k=1}^m \frac{\partial}{\partial R_{ij}} R_{kp} R_{kq} = \partial_{pqij}, \quad \partial_{pqij} = \begin{cases} R_{iq} & \text{if } j = p, p \neq q \\ R_{ip} & \text{if } j = q, p \neq q \\ 2R_{iq} & \text{if } j = p, p = q \\ 0 & \text{otherwise} \end{cases}$$

- (1) Differentiation of Univariate Functions
- (2) Partial Differentiation and Gradients
- (3) Gradients of Vector-Valued Functions
- (4) Gradients of Matrices
- (5) Useful Identities for Computing Gradients
- (6) Backpropagation and Automatic Differentiation
- (7) Higher-Order Derivatives
- (8) Linearization and Multivariate Taylor Series

## Useful Identities

$$\frac{\partial}{\partial \mathbf{X}} \mathbf{f}(\mathbf{X})^\top = \left( \frac{\partial \mathbf{f}(\mathbf{X})}{\partial \mathbf{X}} \right)^\top \quad (5.99)$$

$$\frac{\partial}{\partial \mathbf{X}} \text{tr}(\mathbf{f}(\mathbf{X})) = \text{tr} \left( \frac{\partial \mathbf{f}(\mathbf{X})}{\partial \mathbf{X}} \right) \quad (5.100)$$

$$\frac{\partial}{\partial \mathbf{X}} \det(\mathbf{f}(\mathbf{X})) = \det(\mathbf{f}(\mathbf{X})) \text{tr} \left( \mathbf{f}(\mathbf{X})^{-1} \frac{\partial \mathbf{f}(\mathbf{X})}{\partial \mathbf{X}} \right) \quad (5.101)$$

$$\frac{\partial}{\partial \mathbf{X}} \mathbf{f}(\mathbf{X})^{-1} = -\mathbf{f}(\mathbf{X})^{-1} \frac{\partial \mathbf{f}(\mathbf{X})}{\partial \mathbf{X}} \mathbf{f}(\mathbf{X})^{-1} \quad (5.102)$$

$$\frac{\partial \mathbf{a}^\top \mathbf{X}^{-1} \mathbf{b}}{\partial \mathbf{X}} = -(\mathbf{X}^{-1})^\top \mathbf{a} \mathbf{b}^\top (\mathbf{X}^{-1})^\top \quad (5.103)$$

$$\frac{\partial \mathbf{x}^\top \mathbf{a}}{\partial \mathbf{x}} = \mathbf{a}^\top \quad (5.104)$$

$$\frac{\partial \mathbf{a}^\top \mathbf{x}}{\partial \mathbf{x}} = \mathbf{a}^\top \quad (5.105)$$

$$\frac{\partial \mathbf{a}^\top \mathbf{X} \mathbf{b}}{\partial \mathbf{X}} = \mathbf{a} \mathbf{b}^\top \quad (5.106)$$

$$\frac{\partial \mathbf{x}^\top \mathbf{B} \mathbf{x}}{\partial \mathbf{x}} = \mathbf{x}^\top (\mathbf{B} + \mathbf{B}^\top) \quad (5.107)$$

$$\frac{\partial}{\partial \mathbf{s}} (\mathbf{x} - \mathbf{A}\mathbf{s})^\top \mathbf{W} (\mathbf{x} - \mathbf{A}\mathbf{s}) = -2(\mathbf{x} - \mathbf{A}\mathbf{s})^\top \mathbf{W} \mathbf{A} \quad \text{for symmetric } \mathbf{W} \quad (5.108)$$

- (1) Differentiation of Univariate Functions
- (2) Partial Differentiation and Gradients
- (3) Gradients of Vector-Valued Functions
- (4) Gradients of Matrices
- (5) Useful Identities for Computing Gradients
- (6) **Backpropagation and Automatic Differentiation**
- (7) Higher-Order Derivatives
- (8) Linearization and Multivariate Taylor Series

# Why Should I Care About Backpropagation?

- ▶ **The engine of deep learning:** Without backprop, training neural networks would be **computationally impossible**
- ▶ **Efficiency:** Backprop computes gradients for millions of parameters in roughly the same time as one forward pass
- ▶ **Every AI framework uses it:** PyTorch, TensorFlow, JAX — they all implement automatic differentiation via backprop
- ▶ **Understanding backprop = understanding how AI learns**

**In one sentence:** Backpropagation is *the* algorithm that makes training deep networks practical.

## Motivation: Neural Networks with Many Layers (1)

- ▶ In a neural network with many layers, the function  $\mathbf{y}$  is a many-level function composition

$$\mathbf{y} = (f_K \circ f_{K-1} \circ \cdots \circ f_1)(\mathbf{x}),$$

where, for example,

- ▶  $\mathbf{x}$ : images as inputs,  $\mathbf{y}$ : class labels (e.g., cat or dog) as outputs
- ▶ each  $f_i$  has its own parameters
- ▶ In neural networks, with the model parameters  $\boldsymbol{\theta} = \{\mathbf{A}_0, \mathbf{b}_0, \dots, \mathbf{A}_{K-1}, \mathbf{b}_{K-1}\}$

$$\begin{cases} \mathbf{f}_0 & \stackrel{\text{def}}{=} \mathbf{x} \\ \mathbf{f}_1 & \stackrel{\text{def}}{=} \sigma_1(\mathbf{A}_0 \mathbf{f}_0 + \mathbf{b}_0) \\ \vdots \\ \mathbf{f}_K & \stackrel{\text{def}}{=} \sigma_K(\mathbf{A}_{K-1} \mathbf{f}_{K-1} + \mathbf{b}_{K-1}) \end{cases} \quad \begin{array}{l} \circ \text{ Minimizing the loss function over } \boldsymbol{\theta}: \\ \min_{\boldsymbol{\theta}} L(\boldsymbol{\theta}), \\ \text{ where } L(\boldsymbol{\theta}) = \|\mathbf{y} - \mathbf{f}_K(\boldsymbol{\theta}, \mathbf{x})\|^2 \end{array}$$

- $\sigma_i$  is called the **activation function** at  $i$ -th layer

## Motivation: Neural Networks with Many Layers (2)

- ▶ In neural networks, with the model parameters  $\theta = \{\mathbf{A}_0, \mathbf{b}_0, \dots, \mathbf{A}_{K-1}, \mathbf{b}_{K-1}\}$

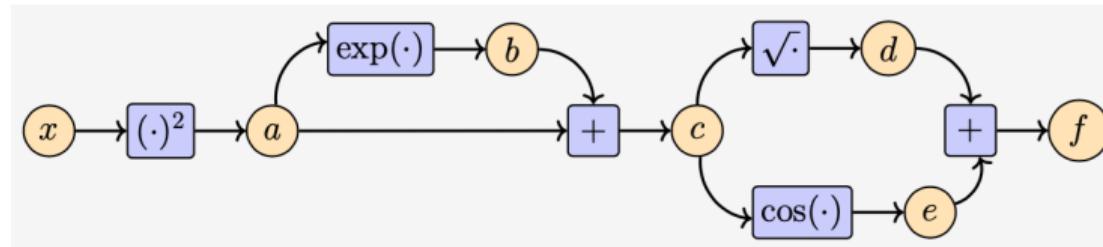
$$\begin{cases} \mathbf{f}_0 & \stackrel{\text{def}}{=} \mathbf{x} \\ \mathbf{f}_1 & \stackrel{\text{def}}{=} \sigma_1(\mathbf{A}_0 \mathbf{f}_0 + \mathbf{b}_0) \\ \vdots \\ \mathbf{f}_K & \stackrel{\text{def}}{=} \sigma_K(\mathbf{A}_{K-1} \mathbf{f}_{K-1} + \mathbf{b}_{K-1}) \end{cases} \quad \begin{array}{l} \circ \text{Minimizing the loss function over } \theta: \\ \min_{\theta} L(\theta), \\ \text{where } L(\theta) = \|\mathbf{y} - \mathbf{f}_K(\theta, \mathbf{x})\|^2 \end{array}$$

- $\sigma_i$  is called the activation function at  $i$ -th layer

- ▶ **Question.** How can we efficiently compute  $\frac{dL}{d\theta}$  in computers?

# Backpropagation: Example (1)

- ▶  $f(x) = \sqrt{x^2 + \exp(x^2)} + \cos(x^2 + \exp(x^2))$
- ▶ Computation graph: Connect via “elementary” operations



$$a = x^2, \quad b = \exp(a), \quad c = a + b, \quad d = \sqrt{c}, \quad e = \cos(c), \quad f = d + e$$

- ▶ Automatic Differentiation
  - ▶ A set of techniques to **numerically** (not symbolically) evaluate the gradient of a function by working with **intermediate variables** and applying the **chain rule**.

## Backpropagation: Example (2)

- ▶  $a = x^2$ ,  $b = \exp(a)$ ,  $c = a + b$ ,  $d = \sqrt{c}$ ,  $e = \cos(c)$ ,  $f = d + e$
- ▶ Derivatives of the intermediate variables with their inputs

$$\frac{\partial a}{\partial x} = 2x, \quad \frac{\partial b}{\partial a} = \exp(a), \quad \frac{\partial c}{\partial a} = 1 = \frac{\partial c}{\partial b}, \quad \frac{\partial d}{\partial c} = \frac{1}{2\sqrt{c}}, \quad \frac{\partial e}{\partial c} = -\sin(c), \quad \frac{\partial f}{\partial d} = 1 = \frac{\partial f}{\partial e}$$

- ▶ Compute  $\frac{\partial f}{\partial x}$  by working backward from the output

$$\frac{\partial f}{\partial c} = \frac{\partial f}{\partial d} \frac{\partial d}{\partial c} + \frac{\partial f}{\partial e} \frac{\partial e}{\partial c}, \quad \frac{\partial f}{\partial b} = \frac{\partial f}{\partial c} \frac{\partial c}{\partial b} \quad \frac{\partial f}{\partial c} = 1 \cdot \frac{1}{2\sqrt{c}} + 1 \cdot (-\sin(c))$$

$$\frac{\partial f}{\partial a} = \frac{\partial f}{\partial b} \frac{\partial b}{\partial a} + \frac{\partial f}{\partial c} \frac{\partial c}{\partial a}, \quad \boxed{\frac{\partial f}{\partial x}} = \frac{\partial f}{\partial a} \frac{\partial a}{\partial x} \quad \frac{\partial f}{\partial b} = \frac{\partial f}{\partial c} \cdot 1, \quad \frac{\partial f}{\partial a} = \frac{\partial f}{\partial b} \exp(a) + \frac{\partial f}{\partial c} \cdot 1$$

$$\boxed{\frac{\partial f}{\partial x}} = \frac{\partial f}{\partial a} \cdot 2x$$

# Quick Recap: Backpropagation

## What we just learned:

- ▶ Build a **computation graph** from elementary operations
- ▶ Compute derivatives of each elementary step
- ▶ Work **backward** from output to input, chaining derivatives
- ▶ Result: efficient gradient of the entire function!
- ▶ **Next up:** Higher-order derivatives and Taylor approximations

# Backpropagation

- ▶ Implementation of gradients can be very expensive, unless we are careful.
- ▶ Using the idea of automatic differentiation, the whole gradient computation is decomposed into a set of gradients of elementary functions and application of the chain rule.
- ▶ Why **backward**?
  - ▶ In neural networks, the input dimensionality is often much higher than the dimensionality of labels.
  - ▶ In this case, backward computation is much cheaper than forward computation.
- ▶ Works if the target is expressed as a computation graph whose elementary functions are differentiable. If not, some care needs to be taken.

- (1) Differentiation of Univariate Functions
- (2) Partial Differentiation and Gradients
- (3) Gradients of Vector-Valued Functions
- (4) Gradients of Matrices
- (5) Useful Identities for Computing Gradients
- (6) Backpropagation and Automatic Differentiation
- (7) Higher-Order Derivatives
- (8) Linearization and Multivariate Taylor Series

# Higher-Order Derivatives

- ▶ Some optimization algorithms (e.g., Newton's method) require second-order derivatives, if they exist.
- ▶ (Truncated) Taylor series is often used as an approximation of a function.
- ▶ For  $f : \mathbb{R}^n \mapsto \mathbb{R}$  of variable  $\mathbf{x} \in \mathbb{R}^n$ ,  $\nabla_{\mathbf{x}} f = \frac{df}{d\mathbf{x}} = \begin{pmatrix} \frac{\partial f(\mathbf{x})}{\partial x_1} & \dots & \frac{\partial f(\mathbf{x})}{\partial x_n} \end{pmatrix} \in \mathbb{R}^{1 \times n}$ 
  - ▶ If  $f$  is twice-differentiable, the order doesn't matter.

$$\nabla_{\mathbf{x}}^2 f = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \vdots & & & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_2 \partial x_n} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{pmatrix}$$

- ▶ For  $f : \mathbb{R}^n \mapsto \mathbb{R}^m$ ,  $\nabla_{\mathbf{x}} f \in \mathbb{R}^{m \times n}$ 
  - ▶ Thus,  $\nabla_{\mathbf{x}}^2 f \in \mathbb{R}^{m \times n \times n}$  (a tensor)

- (1) Differentiation of Univariate Functions
- (2) Partial Differentiation and Gradients
- (3) Gradients of Vector-Valued Functions
- (4) Gradients of Matrices
- (5) Useful Identities for Computing Gradients
- (6) Backpropagation and Automatic Differentiation
- (7) Higher-Order Derivatives
- (8) Linearization and Multivariate Taylor Series

## Function Approximation: Linearization and More

- ▶ First-order approximation of  $f(\mathbf{x})$  (i.e., linearization by taking the first two terms of Taylor Series)

$$f(\mathbf{x}) \approx f(\mathbf{x}_0) + (\nabla_{\mathbf{x}} f)(\mathbf{x}_0)(\mathbf{x} - \mathbf{x}_0)$$

- ▶ Multivariate Taylor Series for  $f : \mathbb{R}^D \mapsto \mathbb{R}$  at  $\mathbf{x}_0$

$$f(\mathbf{x}) = \sum_{k=0}^{\infty} \frac{D_{\mathbf{x}}^k f(\mathbf{x}_0)}{k!} \delta^k,$$

where  $D_{\mathbf{x}}^k f(\mathbf{x}_0)$  is the  $k$ th derivative of  $f$  w.r.t.  $\mathbf{x}$ , evaluated at  $\mathbf{x}_0$ , and  $\delta \stackrel{\text{def}}{=} \mathbf{x} - \mathbf{x}_0$ .

- ▶ Partial sum up to, say  $n$ , can be an approximation of  $f(\mathbf{x})$ .
- ▶  $D_{\mathbf{x}}^k f(\mathbf{x}_0)$  and  $\delta^k$  are  $k$ th order tensors, i.e.,  $k$ -dimensional arrays.
- ▶  $\delta^k$  is a  $k$ -fold outer product  $\otimes$ . For example,  $\delta^2 = \delta \otimes \delta = \delta \delta^\top$ .  $\delta^3 = \delta \otimes \delta \otimes \delta$ .

# Calculus Concepts at a Glance

Concept	Input → Output	Result
Derivative	$f : \mathbb{R} \rightarrow \mathbb{R}$	scalar $f'(x)$
Partial Derivative	$f : \mathbb{R}^n \rightarrow \mathbb{R}$	scalar $\frac{\partial f}{\partial x_i}$
Gradient	$f : \mathbb{R}^n \rightarrow \mathbb{R}$	row vector $\in \mathbb{R}^{1 \times n}$
Jacobian	$\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$	matrix $\in \mathbb{R}^{m \times n}$
Hessian	$f : \mathbb{R}^n \rightarrow \mathbb{R}$	matrix $\in \mathbb{R}^{n \times n}$

**Pattern:** As inputs/outputs grow in dimension, derivatives grow in dimension too!

# Common Mistakes to Avoid

## (1) Forgetting the chain rule in multi-layer functions

Each layer contributes a factor — multiply them all together.

## (2) Confusing gradient shape (row vs. column)

Gradient of  $f : \mathbb{R}^n \rightarrow \mathbb{R}$  is a **row vector** ( $1 \times n$ ), not a column vector.

## (3) Treating partial derivatives as total derivatives

$\frac{\partial f}{\partial x_1}$  holds  $x_2, \dots, x_n$  fixed. The total derivative does not.

## (4) Wrong matrix dimensions in the chain rule

Always check: inner dimensions must match when multiplying Jacobians.

## (5) Forgetting that the Hessian is symmetric

For twice-differentiable  $f$ :  $\frac{\partial^2 f}{\partial x_i \partial x_j} = \frac{\partial^2 f}{\partial x_j \partial x_i}$ .

## Key Takeaways

- (1) Derivative = rate of change; gradient = vector of all partial derivatives
- (2) Chain rule is the foundation of backpropagation: compose simple derivatives to get complex ones
- (3) Jacobian ( $m \times n$  matrix) generalizes the gradient for vector-valued functions  
 $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$
- (4) Backpropagation = efficient gradient computation via chain rule applied backward through a computation graph
- (5) Hessian captures curvature (second-order information) — used in advanced optimizers like Newton's method
- (6) Concept Chain:

Derivative → Gradient → Jacobian → Backprop → Hessian → Taylor

# Questions?

## Review Question 1

- ▶ In machine learning, optimization minimizes a loss function. Explain how the derivative definition

$$\frac{df}{dx} = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

relates to updating model parameters in gradient descent.

## Review Question 2

- ▶ The chain rule is central to deep learning. Explain why

$$(g(f(x)))' = g'(f(x))f'(x)$$

is required to train neural networks with multiple layers.

## Review Question 3

- ▶ For  $f : \mathbb{R}^n \mapsto \mathbb{R}$ , the gradient gives the direction of steepest ascent. Explain how gradients are used to update model parameters during training.

## Review Question 4

- ▶ Consider a neural network layer written as

$$\mathbf{f}(\mathbf{x}) = \mathbf{Ax}.$$

Show that its Jacobian is  $\mathbf{A}$ . What does this represent in terms of feature transformation in deep learning?

## Review Question 5

- ▶ In computer vision, images are represented as high-dimensional vectors. Explain how partial derivatives measure how sensitive the loss is to small changes in pixel values.

## Review Question 6

- ▶ Using the chain rule,

$$\frac{df}{dt} = \frac{\partial f}{\partial x_1} \frac{\partial x_1}{\partial t} + \frac{\partial f}{\partial x_2} \frac{\partial x_2}{\partial t},$$

explain how backpropagation computes gradients through multiple layers.

## Review Question 7

- ▶ For the least-squares loss

$$L(\theta) = \|\mathbf{y} - \Phi\theta\|^2,$$

explain what the matrix  $\Phi$  represents in machine learning and how minimizing this loss corresponds to fitting a model.

## Review Question 8

- ▶ Explain why backpropagation is computationally efficient when training deep neural networks with many parameters.

## Review Question 9

- ▶ The Hessian matrix contains second-order derivatives. Explain how curvature information can influence optimization in machine learning (e.g., faster convergence or instability).

## Review Question 10

- ▶ Taylor expansion gives a local approximation:

$$f(\mathbf{x}) \approx f(\mathbf{x}_0) + (\nabla f)(\mathbf{x}_0)(\mathbf{x} - \mathbf{x}_0).$$

Explain how this idea relates to linearization in machine learning models and local behavior of neural networks.

# Theory (1): Backprop as Vector–Jacobian Products (VJP)

- ▶ Consider a composition

$$\mathbf{f}_K = (f_K \circ f_{K-1} \circ \cdots \circ f_1)(\mathbf{x}), \quad L = \ell(\mathbf{f}_K).$$

- ▶ Jacobians:

$$\mathbf{J}_i \stackrel{\text{def}}{=} \frac{df_i}{d\mathbf{f}_{i-1}} \quad \Rightarrow \quad \frac{d\mathbf{f}_K}{d\mathbf{x}} = \mathbf{J}_K \mathbf{J}_{K-1} \cdots \mathbf{J}_1.$$

- ▶ Naively storing  $\mathbf{J}_i$  is infeasible. Backprop computes VJP:

$$\underbrace{\frac{dL}{d\mathbf{f}_{i-1}}}_{1 \times d_{i-1}} = \underbrace{\frac{dL}{d\mathbf{f}_i}}_{1 \times d_i} \underbrace{\mathbf{J}_i}_{d_i \times d_{i-1}}.$$

- ▶ Practical takeaway:

- ▶ Backward mode is efficient when output is scalar (typical loss).
- ▶ Complexity is comparable to forward pass (up to constant factors).

## Theory (2): Hessian Structure and Curvature-Aware Updates

- ▶ For parameters  $\theta \in \mathbb{R}^P$ , gradient and Hessian:

$$\nabla L(\theta) \in \mathbb{R}^{1 \times P}, \quad \nabla^2 L(\theta) \in \mathbb{R}^{P \times P}.$$

- ▶ Second-order Taylor expansion:

$$L(\theta + \Delta) \approx L(\theta) + \nabla L(\theta)\Delta + \frac{1}{2}\Delta^\top \nabla^2 L(\theta)\Delta.$$

- ▶ Newton step (ideal):

$$\Delta^* = -\nabla^2 L(\theta)^{-1} \nabla L(\theta)^\top.$$

- ▶ Why we care in deep nets:

- ▶ Sharp directions (large eigenvalues of  $\nabla^2$ ) constrain step size.
- ▶ Flat directions (small eigenvalues) dominate generalization behavior.

- ▶ Common approximation idea (diagonal / low-rank):

$$\nabla^2 \approx \mathbf{U}_k \Lambda_k \mathbf{U}_k^\top + \lambda \mathbf{I}.$$

## Theory (3): Gauss–Newton / Fisher from Jacobians

- ▶ For squared loss with residual  $\mathbf{r}(\theta) \in \mathbb{R}^N$ :

$$L(\theta) = \frac{1}{2} \|\mathbf{r}(\theta)\|^2, \quad \mathbf{r} = \mathbf{y} - \hat{\mathbf{y}}(\theta).$$

- ▶ Jacobian of residuals:

$$\mathbf{J}_r \stackrel{\text{def}}{=} \frac{d\mathbf{r}}{d\theta} \in \mathbb{R}^{N \times P}.$$

- ▶ Gradient:

$$\nabla L(\theta) = \mathbf{r}^\top \mathbf{J}_r.$$

- ▶ Exact Hessian:

$$\nabla^2 L(\theta) = \mathbf{J}_r^\top \mathbf{J}_r + \sum_{i=1}^N r_i \nabla^2 r_i.$$

- ▶ **Gauss–Newton approximation** (drop the second term):

$$\nabla^2 L(\theta) \approx \mathbf{J}_r^\top \mathbf{J}_r,$$

which is PSD and captures curvature near good fits (small  $r_i$ ).

- ▶ Link to Fisher (probabilistic view): curvature  $\approx$  covariance of score.

## Theory (4): Linearization and Adversarial Sensitivity

- ▶ First-order change in loss under input perturbation  $\delta$ :

$$L(\mathbf{x} + \delta) \approx L(\mathbf{x}) + \nabla_{\mathbf{x}} L(\mathbf{x}) \delta.$$

- ▶ Worst-case (under  $\ell_2$ -budget  $\|\delta\|_2 \leq \epsilon$ ):

$$\max_{\|\delta\|_2 \leq \epsilon} \nabla_{\mathbf{x}} L \delta = \epsilon \|\nabla_{\mathbf{x}} L\|_2.$$

- ▶ For  $\ell_\infty$ -budget  $\|\delta\|_\infty \leq \epsilon$ :

$$\delta^* = \epsilon \operatorname{sign}(\nabla_{\mathbf{x}} L(\mathbf{x})) \quad (\text{FGSM-type direction}).$$

- ▶ Vision connection:

- ▶ Robustness is controlled by **Jacobian norms** of the network.
- ▶ Regularizers such as  $\|\mathbf{J}_{\mathbf{x}}\|$  penalize local sensitivity.

## Theory (5): Jacobian Rank and Feature Collapse

- ▶ Let  $\mathbf{f}(\mathbf{x}) \in \mathbb{R}^d$  be a learned representation (penultimate layer). Jacobian:

$$\mathbf{J}_{\mathbf{x}} \stackrel{\text{def}}{=} \frac{d\mathbf{f}}{d\mathbf{x}} \in \mathbb{R}^{d \times D}.$$

- ▶ Local dimension is governed by  $\text{rank}(\mathbf{J}_{\mathbf{x}})$ :

- ▶ If  $\text{rank}(\mathbf{J}_{\mathbf{x}}) \ll \min(d, D)$ , features vary only in a low-dim subspace.
- ▶ This is a differential view of **representation collapse**.

- ▶ A useful summary:

$$\text{effective dim}(\mathbf{x}) \approx \text{rank}(\mathbf{J}_{\mathbf{x}}) \quad \text{or} \quad \text{tr}(\mathbf{J}_{\mathbf{x}} \mathbf{J}_{\mathbf{x}}^\top).$$

- ▶ Practical link:

- ▶ Self-supervised methods control collapse by constraints on covariance / spectrum.
- ▶ In CNNs, collapse is often layerwise and spatially structured.

# PhD Level Content Only

The next slides are **research-level** vector calculus for deep learning theory:

- ▶ Implicit bias via gradient flow
- ▶ NTK / kernel regression dynamics
- ▶ Fisher geometry and natural gradients
- ▶ Spectral structure of Hessians in vision models
- ▶ Differential view of invariances (Lie groups)

Not intended for undergraduate coursework

## PhD (6): Gradient Flow and Implicit Regularization

- ▶ Continuous-time limit of gradient descent:

$$\frac{d\theta(t)}{dt} = -\nabla_{\theta} L(\theta(t))^{\top}.$$

- ▶ For linear regression with separable data and certain losses, gradient flow selects a **minimum-norm** solution (implicit bias).
- ▶ Prototype statement (informal but useful):

$\theta(t)$  converges to  $\arg \min_{\theta} \|\theta\|_2$  s.t. zero training error.

- ▶ In deep linear networks, the induced implicit bias often corresponds to **low-rank** / **nuclear-norm-like** structure on end-to-end maps.
- ▶ Conceptual bridge to your Lecture 4:

Training  $\Rightarrow$  spectral shrinkage of singular values.

## PhD (7): NTK as a Spectral Learning Operator

- ▶ Linearization at initialization  $\theta_0$ :

$$f_{\theta}(x) \approx f_{\theta_0}(x) + \nabla_{\theta} f_{\theta_0}(x)^{\top} (\theta - \theta_0).$$

- ▶ Define NTK kernel on dataset  $\{x_i\}$ :

$$\mathbf{K}_{ij} = \langle \nabla_{\theta} f(x_i), \nabla_{\theta} f(x_j) \rangle.$$

- ▶ Gradient descent on squared loss yields (in function space):

$$\hat{\mathbf{y}}(t) = \mathbf{y} - e^{-\eta \mathbf{K} t} (\mathbf{y} - \hat{\mathbf{y}}(0)).$$

- ▶ Spectral decomposition  $\mathbf{K} = \mathbf{U} \Lambda \mathbf{U}^{\top}$  gives modewise rates:

mode  $k$  converges at rate  $e^{-\eta \lambda_k t}$ .

- ▶ Vision implication: **spectrum of  $\mathbf{K}$**  determines which patterns are learned first.

## PhD (8): Fisher Information and Natural Gradient

- ▶ Probabilistic model  $p(y|x, \theta)$  with negative log-likelihood loss  $L(\theta)$ .
- ▶ Fisher information:

$$\mathbf{F}(\theta) = \mathbb{E}\left[\nabla_{\theta} \log p(y|x, \theta) \nabla_{\theta} \log p(y|x, \theta)^{\top}\right].$$

- ▶ Natural gradient step (Riemannian metric induced by  $\mathbf{F}$ ):

$$\theta^+ = \theta - \eta \mathbf{F}(\theta)^{-1} \nabla_{\theta} L(\theta)^{\top}.$$

- ▶ Spectral view:

$$\mathbf{F} = \mathbf{U} \Lambda \mathbf{U}^{\top} \Rightarrow \mathbf{F}^{-1} = \mathbf{U} \Lambda^{-1} \mathbf{U}^{\top}.$$

- ▶ Interpretation:
  - ▶ learning is slowed along **high-curvature** directions
  - ▶ preconditioning equalizes progress across parameter manifolds

## PhD (9): Hessian Spectrum, Sharpness, and Generalization

- ▶ Local quadratic model:

$$L(\theta + \Delta) \approx L(\theta) + \nabla L \Delta + \frac{1}{2} \Delta^\top \nabla^2 L \Delta.$$

- ▶ Let  $\nabla^2 L = \mathbf{U} \Lambda \mathbf{U}^\top$  with eigenvalues  $\lambda_1 \geq \dots \geq \lambda_P$ .
- ▶ Sharpness proxy (spectral):

$$\text{sharpness} \sim \lambda_1 \quad \text{or} \quad \text{tr}(\nabla^2 L) = \sum_i \lambda_i.$$

- ▶ In large vision nets, empirical Hessians are often:
  - ▶ **low effective rank** (few large eigenvalues)
  - ▶ heavy-tailed spectra (many small but nontrivial modes)
- ▶ Practical bridge:

SVD / low-rank approx  $\Rightarrow$  efficient curvature modeling.

# PhD (10): Differential Invariances via Lie Groups

- ▶ Suppose a transformation group  $g(\alpha)$  acts on inputs:  $x \mapsto g(\alpha) \cdot x$ .
- ▶ Invariance objective:

$$f(g(\alpha) \cdot x) \approx f(x) \quad \forall \alpha \text{ small.}$$

- ▶ Differentiate at identity ( $\alpha = 0$ ):

$$\frac{d}{d\alpha} f(g(\alpha) \cdot x) \Big|_{\alpha=0} = \nabla_x f(x) \frac{d}{d\alpha} (g(\alpha) \cdot x) \Big|_{\alpha=0} \approx 0.$$

- ▶ This yields a constraint on the gradient:

$$\nabla_x f(x) \mathbf{T}(x) \approx 0,$$

where  $\mathbf{T}(x)$  spans the tangent space of the orbit (translations, rotations, etc.).

- ▶ Vision takeaway:
  - ▶ CNN inductive bias approximates invariance by architectural constraints.
  - ▶ Gradient-based regularizers can enforce invariances at training time.