

Roadmap

- (1) Determinant and Trace
- (2) Eigenvalues and Eigenvectors
- (3) Cholesky Decomposition
- (4) Eigendecomposition and Diagonalization
- (5) Singular Value Decomposition
- (6) Matrix Approximation
- (7) Matrix Phylogeny

- ▶ How to summarize matrices: determinants and eigenvalues
- ▶ How matrices can be decomposed: Cholesky decomposition, diagonalization, singular value decomposition
- ▶ How these decompositions can be used for matrix approximation

Roadmap

- (1) Determinant and Trace
- (2) Eigenvalues and Eigenvectors
- (3) Cholesky Decomposition
- (4) Eigendecomposition and Diagonalization
- (5) Singular Value Decomposition
- (6) Matrix Approximation
- (7) Matrix Phylogeny

Determinant: Motivation (1)

- ▶ For $\mathbf{A} = \begin{pmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{pmatrix}$, $\mathbf{A}^{-1} = \frac{1}{a_{11}a_{22} - a_{12}a_{21}} \begin{pmatrix} a_{22} & -a_{12} \\ -a_{21} & a_{11} \end{pmatrix}$.
- ▶ \mathbf{A} is invertible iff $a_{11}a_{22} - a_{12}a_{21} \neq 0$
- ▶ Let's define $\det(\mathbf{A}) = a_{11}a_{22} - a_{12}a_{21}$.
- ▶ Notation: $\det(\mathbf{A})$ or $|\text{whole matrix}|$
- ▶ What about 3×3 matrix? By doing some algebra (e.g., Gaussian elimination),

$$\begin{vmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{vmatrix} = a_{11}a_{22}a_{33} + a_{21}a_{32}a_{13} + a_{31}a_{12}a_{23} - a_{31}a_{22}a_{13} - a_{11}a_{32}a_{23} - a_{21}a_{12}a_{33}$$

Determinant: Motivation (2)

- ▶ Try to find some pattern ...

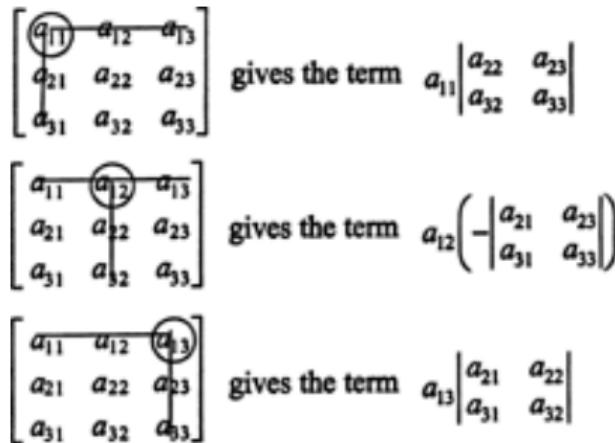
$$a_{11}a_{22}a_{33} + a_{21}a_{32}a_{13} + a_{31}a_{12}a_{23}$$

$$- a_{31}a_{22}a_{13} - a_{11}a_{32}a_{23} - a_{21}a_{12}a_{33} =$$

$$\begin{aligned} & a_{11}(-1)^{1+1} \det(\mathbf{A}_{1,1}) + a_{12}(-1)^{1+2} \det(\mathbf{A}_{1,2}) \\ & + a_{13}(-1)^{1+3} \det(\mathbf{A}_{1,3}) \end{aligned}$$

- $\mathbf{A}_{k,j}$ is the submatrix of \mathbf{A} that we obtain when deleting row k and column j .

- ▶ This is called **Laplace expansion**.
- ▶ Now, we can generalize this and provide the formal definition of determinant.



source: www.cliffsnotes.com

Determinant: Formal Definition

Determinant

For a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$, for all $j = 1, \dots, n$,

$$\text{Expansion along column } j: \det(\mathbf{A}) = \sum_{k=1}^n (-1)^{k+j} a_{kj} \det(\mathbf{A}_{k,j})$$

$$\text{Expansion along row } j: \det(\mathbf{A}) = \sum_{k=1}^n (-1)^{k+j} a_{jk} \det(\mathbf{A}_{j,k})$$

- ▶ All expansions are equal, so no problem with the definition.
- ▶ **Theorem.** $\det(\mathbf{A}) \neq 0 \iff \text{rank}(\mathbf{A}) = n \iff \mathbf{A} \text{ is invertible.}$

Determinant: Properties

- (1) $\det(\mathbf{AB}) = \det(\mathbf{A}) \det(\mathbf{B})$
- (2) $\det(\mathbf{A}) = \det(\mathbf{A}^\top)$
- (3) For a regular \mathbf{A} , $\det(\mathbf{A}^{-1}) = 1 / \det(\mathbf{A})$
- (4) For two similar matrices \mathbf{A}, \mathbf{A}' (i.e., $\mathbf{A}' = \mathbf{S}^{-1}\mathbf{AS}$ for some \mathbf{S}), $\det(\mathbf{A}) = \det(\mathbf{A}')$
- (5) For a triangular matrix¹ \mathbf{T} , $\det(\mathbf{T}) = \prod_{i=1}^n T_{ii}$
- (6) Adding a multiple of a column/row to another one does not change $\det(\mathbf{A})$
- (7) Multiplication of a column/row with λ scales det by λ . Hence,
 $\det(\lambda\mathbf{A}) = \lambda^n \det(\mathbf{A})$
- (8) Swapping two rows/columns changes the sign of $\det(\mathbf{A})$
 - o Using (5)-(8), Gaussian elimination (reaching a triangular matrix) enables computing the determinant.

¹This includes diagonal matrices.

- ▶ **Definition.** The trace of a square matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ is defined as

$$\text{tr}(\mathbf{A}) \stackrel{\text{def}}{=} \sum_{i=1}^n a_{ii}$$

- ▶ $\text{tr}(\mathbf{A} + \mathbf{B}) = \text{tr}(\mathbf{A}) + \text{tr}(\mathbf{B})$
- ▶ $\text{tr}(\alpha\mathbf{A}) = \alpha\text{tr}(\mathbf{A})$
- ▶ $\text{tr}(\mathbf{I}_n) = n$

Invariant under Cyclic Permutations

- ▶ $\text{tr}(\mathbf{AB}) = \text{tr}(\mathbf{BA})$ for $\mathbf{A} \in \mathbb{R}^{n \times k}$ and $\mathbf{B} \in \mathbb{R}^{k \times n}$
- ▶ $\text{tr}(\mathbf{AKL}) = \text{tr}(\mathbf{KLA})$, for $\mathbf{A} \in \mathbb{R}^{a \times k}$, $\mathbf{K} \in \mathbb{R}^{k \times l}$, $\mathbf{L} \in \mathbb{R}^{l \times a}$
- ▶ $\text{tr}(\mathbf{xy}^\top) = \text{tr}(\mathbf{y}^\top \mathbf{x}) = \mathbf{y}^\top \mathbf{x} \in \mathbb{R}$
- ▶ A linear mapping $\Phi : V \rightarrow V$ can be represented by a matrix \mathbf{A} or another matrix \mathbf{B} .
 - ▶ \mathbf{A} and \mathbf{B} use different bases, where $\mathbf{B} = \mathbf{S}^{-1}\mathbf{AS}$

$$\text{tr}(\mathbf{B}) = \text{tr}(\mathbf{S}^{-1}\mathbf{AS}) = \text{tr}(\mathbf{ASS}^{-1}) = \text{tr}(\mathbf{A})$$

- ▶ **Message.** While matrix representations of linear mappings are basis dependent, their traces are not.

Background: Characteristic Polynomial

- **Definition.** For $\lambda \in \mathbb{R}$ and a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$, the characteristic polynomial of \mathbf{A} is defined as:

$$\begin{aligned} p_{\mathbf{A}}(\lambda) &\stackrel{\text{def}}{=} \det(\mathbf{A} - \lambda \mathbf{I}) \\ &= c_0 + c_1 \lambda + c_2 \lambda^2 + \cdots + c_{n-1} \lambda^{n-1} + (-1)^n \lambda^n, \end{aligned}$$

where $c_0 = \det(\mathbf{A})$ and $c_{n-1} = (-1)^{n-1} \text{tr}(\mathbf{A})$.

- **Example.** For $\mathbf{A} = \begin{pmatrix} 4 & 2 \\ 1 & 3 \end{pmatrix}$,

$$p_{\mathbf{A}}(\lambda) = \begin{vmatrix} 4 - \lambda & 2 \\ 1 & 3 - \lambda \end{vmatrix} = (4 - \lambda)(3 - \lambda) - 2 \cdot 1$$

What we just learned:

- ▶ **Determinant** = one number that tells you if a matrix is invertible ($\det(\mathbf{A}) \neq 0$)
- ▶ **Trace** = sum of diagonal entries
- ▶ Both are **invariant** under basis change — they describe the matrix itself, not a particular representation
- ▶ **Next up:** Eigenvalues and eigenvectors — the “DNA” of a matrix

- (1) Determinant and Trace
- (2) Eigenvalues and Eigenvectors
- (3) Cholesky Decomposition
- (4) Eigendecomposition and Diagonalization
- (5) Singular Value Decomposition
- (6) Matrix Approximation
- (7) Matrix Phylogeny

Why Should I Care About Eigenvalues?

- ▶ **Google's PageRank:** Ranks billions of web pages using the dominant eigenvector of the web link matrix
- ▶ **Vibration analysis:** Engineers find eigenvalues to predict when a bridge might resonate and collapse
- ▶ **Facial recognition:** “Eigenfaces” — the most important patterns in face images are eigenvectors
- ▶ **Machine Learning:** PCA (Principal Component Analysis) uses eigenvalues to reduce data from 1000s of features to just a few

In one sentence: Eigenvalues reveal the most important directions hidden inside data.

Eigenvalue and Eigenvector

- ▶ **Definition.** Consider a square matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$. Then, $\lambda \in \mathbb{R}$ is an eigenvalue of \mathbf{A} and $\mathbf{x} \in \mathbb{R}^n \setminus \{0\}$ is the corresponding eigenvector of \mathbf{A} if

$$\mathbf{A} \mathbf{x} = \lambda \mathbf{x}$$

- ▶ Equivalent statements
 - ▶ λ is an eigenvalue.
 - ▶ $(\mathbf{A} - \lambda \mathbf{I}_n) \mathbf{x} = 0$ can be solved non-trivially, i.e., $\mathbf{x} \neq \mathbf{0}$.
 - ▶ $\text{rank}(\mathbf{A} - \lambda \mathbf{I}_n) < n$.
 - ▶ $\det(\mathbf{A} - \lambda \mathbf{I}_n) = 0 \iff$ The characteristic polynomial $p_{\mathbf{A}}(\lambda) = 0$.

Example

- ▶ For $\mathbf{A} = \begin{pmatrix} 4 & 2 \\ 1 & 3 \end{pmatrix}$, $p_{\mathbf{A}}(\lambda) = \begin{vmatrix} 4 - \lambda & 2 \\ 1 & 3 - \lambda \end{vmatrix} = (4 - \lambda)(3 - \lambda) - 2 \cdot 1 = \lambda^2 - 7\lambda + 10$
- ▶ Eigenvalues $\lambda = 2$ or $\lambda = 5$.
- ▶ Eigenspace E_5 for $\lambda = 5$

$$\begin{pmatrix} 4 - \lambda & 2 \\ 1 & 3 - \lambda \end{pmatrix} \mathbf{x} = 0 \implies \begin{pmatrix} -1 & 2 \\ 1 & -2 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = 0 \implies E_5 = \text{span} \begin{pmatrix} 2 \\ 1 \end{pmatrix}$$

- ▶ Eigenspace E_2 for $\lambda = 2$. Similarly, we get $E_2 = \text{span} \begin{pmatrix} 1 \\ -1 \end{pmatrix}$
- ▶ **Message.** Eigenvectors are not unique.

Properties (1)

- ▶ If \mathbf{x} is an eigenvector of \mathbf{A} , so are all vectors that are collinear².
- ▶ E_λ : the set of all eigenvectors for eigenvalue λ , spanning a subspace of \mathbb{R}^n . We call this the **eigenspace** of \mathbf{A} for λ .
- ▶ E_λ is the solution space of $(\mathbf{A} - \lambda\mathbf{I})\mathbf{x} = 0$, thus $E_\lambda = \ker(\mathbf{A} - \lambda\mathbf{I})$
- ▶ **Geometric interpretation**
 - ▶ The eigenvector corresponding to a nonzero eigenvalue points in a direction stretched by the linear mapping.
 - ▶ The eigenvalue is the factor of stretching.
- ▶ Identity matrix \mathbf{I} : one eigenvalue $\lambda = 1$ and all vectors $\mathbf{x} \neq \mathbf{0}$ are eigenvectors.

²Two vectors are collinear if they point in the same or the opposite direction.

Properties (2)

- ▶ \mathbf{A} and \mathbf{A}^\top share the eigenvalues, but not necessarily eigenvectors.
- ▶ For two similar matrices \mathbf{A}, \mathbf{A}' (i.e., $\mathbf{A}' = \mathbf{S}^{-1}\mathbf{AS}$ for some \mathbf{S}), they possess the same eigenvalues.
 - ▶ Meaning: A linear mapping Φ has eigenvalues that are **independent** of the choice of basis of its transformation matrix.
 - ▶ Symmetric, positive definite matrices always have **positive, real** eigenvalues.

determinant, trace, eigenvalues: all **invariant** under basis change

Examples for Geometric Interpretation (1)

$$\mathbf{A} = \begin{pmatrix} \frac{1}{2} & 0 \\ 0 & 2 \end{pmatrix}, \det(\mathbf{A}) = 1$$

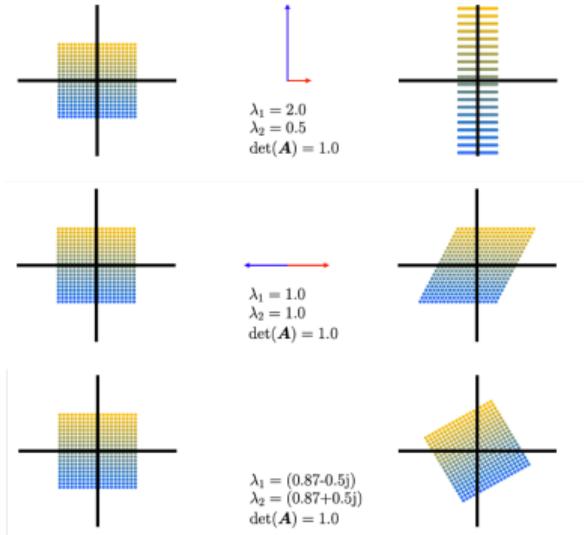
- ▶ $\lambda_1 = \frac{1}{2}, \lambda_2 = 2$
- ▶ eigenvectors: canonical basis vectors
- ▶ area preserving, just vertical and horizontal stretching.

$$\mathbf{A} = \begin{pmatrix} 1 & \frac{1}{2} \\ 0 & 1 \end{pmatrix}, \det(\mathbf{A}) = 1$$

- ▶ $\lambda_1 = \lambda_2 = 1$
- ▶ eigenvectors: collinear over the horizontal line
- ▶ area preserving, shearing

$$\mathbf{A} = \begin{pmatrix} \cos\left(\frac{\pi}{6}\right) & -\sin\left(\frac{\pi}{6}\right) \\ \sin\left(\frac{\pi}{6}\right) & \cos\left(\frac{\pi}{6}\right) \end{pmatrix}, \det(\mathbf{A}) = 1$$

- ▶ Rotation by $\pi/6$ counter-clockwise
- ▶ only complex eigenvalues (no real eigenvectors)
- ▶ area preserving



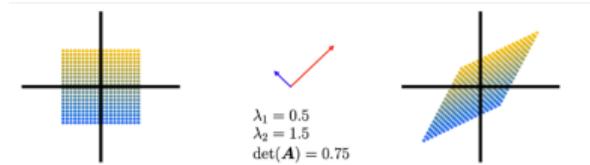
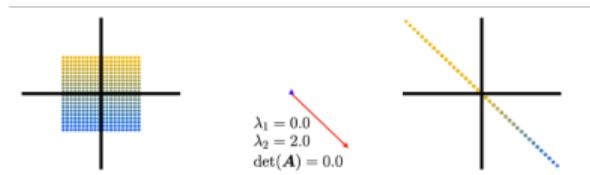
Examples for Geometric Interpretation (2)

4. $\mathbf{A} = \begin{pmatrix} 1 & -1 \\ -1 & 1 \end{pmatrix}$, $\det(\mathbf{A}) = 0$

- ▶ $\lambda_1 = 0, \lambda_2 = 2$
- ▶ Mapping that collapses a 2D onto 1D
- ▶ area collapses

5. $\mathbf{A} = \begin{pmatrix} 1 & \frac{1}{2} \\ \frac{1}{2} & 1 \end{pmatrix}$, $\det(\mathbf{A}) = 3/4$

- ▶ $\lambda_1 = 0.5, \lambda_2 = 1.5$
- ▶ area scales by 75%, shearing and stretching



Properties (3)

- ▶ For $\mathbf{A} \in \mathbb{R}^{n \times n}$, n distinct eigenvalues \implies eigenvectors are linearly independent, which form a basis of \mathbb{R}^n .
 - ▶ Converse is not true.
 - ▶ Example of n linearly independent eigenvectors for fewer than n distinct eigenvalues?
- ▶ **Determinant.** For (possibly repeated) eigenvalues λ_i of $\mathbf{A} \in \mathbb{R}^{n \times n}$,

$$\det(\mathbf{A}) = \prod_{i=1}^n \lambda_i$$

- ▶ **Trace.** For (possibly repeated) eigenvalues λ_i of $\mathbf{A} \in \mathbb{R}^{n \times n}$,

$$\text{tr}(\mathbf{A}) = \sum_{i=1}^n \lambda_i$$

- ▶ **Message.** $\det(\mathbf{A})$ is the **area scaling** and $\text{tr}(\mathbf{A})$ is the **circumference scaling**

Quick Recap: Eigenvalues & Eigenvectors

What we just learned:

- ▶ An **eigenvector** is a direction that a matrix only **stretches** (not rotates)
- ▶ The **eigenvalue** tells you **how much** it stretches
- ▶ $\mathbf{Ax} = \lambda\mathbf{x}$ — the matrix acts like simple multiplication along special directions
- ▶ $\det(\mathbf{A}) = \prod \lambda_i$ and $\text{tr}(\mathbf{A}) = \sum \lambda_i$
- ▶ **Next up:** Cholesky decomposition — a matrix “square root”

Roadmap

- (1) Determinant and Trace
- (2) Eigenvalues and Eigenvectors
- (3) Cholesky Decomposition
- (4) Eigendecomposition and Diagonalization
- (5) Singular Value Decomposition
- (6) Matrix Approximation
- (7) Matrix Phylogeny

Why Should I Care About Cholesky?

- ▶ **Speed:** Solving $\mathbf{Ax} = \mathbf{b}$ with Cholesky is **twice as fast** as general methods
- ▶ **Machine Learning:** Every Gaussian distribution has a covariance matrix — Cholesky decomposes it efficiently
- ▶ **Numerical stability:** More stable than LU for symmetric positive definite matrices
- ▶ **Sampling:** To generate random data that follows a specific correlation pattern, you need Cholesky

In one sentence: Cholesky is the fast, stable way to work with symmetric positive definite matrices.

LU Decomposition

$$\begin{bmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ 0 & & \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{bmatrix}$$

Upper Triangular Matrix

$$\begin{bmatrix} \cdot & & & 0 \\ \cdot & \cdot & & \\ \cdot & \cdot & \cdot & \\ \cdot & \cdot & \cdot & \\ \cdot & \cdot & \cdot & \cdot \end{bmatrix}$$

Lower Triangular Matrix

Source: <http://mathonline.wikidot.com/>

- ▶ Gaussian elimination transforms \mathbf{A} into an upper triangular matrix \mathbf{U} .
- ▶ Each elimination step is an elementary matrix \mathbf{E}_i acting on the left:

$$\mathbf{E}_k \mathbf{E}_{k-1} \cdots \mathbf{E}_1 \mathbf{A} = \mathbf{U}.$$

- ▶ Hence,

$$\mathbf{A} = (\mathbf{E}_1^{-1} \mathbf{E}_2^{-1} \cdots \mathbf{E}_k^{-1}) \mathbf{U} = \mathbf{L} \mathbf{U}.$$

- ▶ The matrices \mathbf{E}_i are lower triangular; so are their inverses and their product \mathbf{L} .

Cholesky Decomposition

- ▶ Analogy for real numbers: a square root, e.g., $9 = 3 \times 3$
- ▶ **Theorem.** For a symmetric, positive definite matrix \mathbf{A} , $\mathbf{A} = \mathbf{L}\mathbf{L}^\top$, where
 - ▶ \mathbf{L} is a lower-triangular matrix with positive diagonals
 - ▶ Such a \mathbf{L} is unique, called **Cholesky factor** of \mathbf{A} .
- ▶ Applications
 - (a) factorization of covariance matrix of a multivariate Gaussian variable
 - (b) linear transformation of random variables
 - (c) fast determinant computation: $\det(\mathbf{A}) = \det(\mathbf{L})\det(\mathbf{L}^\top) = \det(\mathbf{L})^2$, where $\det(\mathbf{L}) = \prod_i l_{ii}$. Thus, $\det(\mathbf{A}) = \prod_i l_{ii}^2$.

Quick Recap: Cholesky Decomposition

What we just learned:

- ▶ **LU:** Any square matrix = lower \times upper triangular
- ▶ **Cholesky:** For symmetric positive definite matrices: $\mathbf{A} = \mathbf{LL}^T$
- ▶ Think of it as the “square root” of a matrix
- ▶ **Next up:** Eigendecomposition — breaking a matrix into its fundamental pieces

Roadmap

- (1) Determinant and Trace
- (2) Eigenvalues and Eigenvectors
- (3) Cholesky Decomposition
- (4) Eigendecomposition and Diagonalization
- (5) Singular Value Decomposition
- (6) Matrix Approximation
- (7) Matrix Phylogeny

Why Should I Care About Eigendecomposition?

- ▶ **Powers of matrices:** Instead of multiplying $\mathbf{A} \times \mathbf{A} \times \cdots \times \mathbf{A}$, just compute $\mathbf{P}\mathbf{D}^k\mathbf{P}^{-1}$ — exponentially faster!
- ▶ **Understanding transformations:** Eigendecomposition shows you exactly *what* a matrix does: stretch along specific directions
- ▶ **Stability analysis:** Are eigenvalues < 1 ? System is stable. > 1 ? System explodes.
- ▶ **Quantum mechanics:** Observable quantities (energy, momentum) are eigenvalues of operators

In one sentence: Eigendecomposition breaks a matrix into its simplest building blocks.

Diagonal Matrix and Diagonalization

- **Diagonal matrix.** zero on all off-diagonal elements, $\mathbf{D} = \begin{pmatrix} d_1 & \cdots & 0 \\ \vdots & & \vdots \\ 0 & \cdots & d_n \end{pmatrix}$
- $\mathbf{D}^k = \begin{pmatrix} d_1^k & \cdots & 0 \\ \vdots & & \vdots \\ 0 & \cdots & d_n^k \end{pmatrix}, \quad \mathbf{D}^{-1} = \begin{pmatrix} 1/d_1 & \cdots & 0 \\ \vdots & & \vdots \\ 0 & \cdots & 1/d_n \end{pmatrix}, \quad \det(\mathbf{D}) = d_1 d_2 \cdots d_n$
- **Definition.** $\mathbf{A} \in \mathbb{R}^{n \times n}$ is **diagonalizable** if it is similar to a diagonal matrix \mathbf{D} , i.e., \exists an **invertible** $\mathbf{P} \in \mathbb{R}^{n \times n}$, such that $\mathbf{D} = \mathbf{P}^{-1}\mathbf{A}\mathbf{P}$.
- **Definition.** $\mathbf{A} \in \mathbb{R}^{n \times n}$ is **orthogonally diagonalizable** if it is similar to a diagonal matrix \mathbf{D} , i.e., \exists an **orthogonal** $\mathbf{P} \in \mathbb{R}^{n \times n}$, such that $\mathbf{D} = \mathbf{P}^{-1}\mathbf{A}\mathbf{P} = \mathbf{P}^\top\mathbf{A}\mathbf{P}$.

Power of Diagonalization

- ▶ $\mathbf{A}^k = \mathbf{P}\mathbf{D}^k\mathbf{P}^{-1}$
- ▶ $\det(\mathbf{A}) = \det(\mathbf{P}) \det(\mathbf{D}) \det(\mathbf{P}^{-1}) = \det(\mathbf{D}) = \prod_i d_{ii}$
- ▶ Many other things ...
- ▶ **Question.** Under what condition is \mathbf{A} diagonalizable (or orthogonally diagonalizable) and how can we find \mathbf{P} (thus \mathbf{D})?

Diagonalizability, Algebraic/Geometric Multiplicity

- ▶ **Definition.** For a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ with an eigenvalue λ_i ,
 - ▶ the **algebraic multiplicity** α_i of λ_i is the number of times the root appears in the characteristic polynomial.
 - ▶ the **geometric multiplicity** ζ_i of λ_i is the number of linearly independent eigenvectors associated with λ_i (i.e., the dimension of the eigenspace spanned by the eigenvectors of λ_i)
- ▶ **Example.** The matrix $\mathbf{A} = \begin{pmatrix} 2 & 1 \\ 0 & 2 \end{pmatrix}$ has two repeated eigenvalues $\lambda_1 = \lambda_2 = 2$, thus $\alpha_1 = 2$. However, it has only one distinct unit eigenvector $\mathbf{x} = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$, thus $\zeta_1 = 1$.
- ▶ **Theorem.** $\mathbf{A} \in \mathbb{R}^{n \times n}$ is **diagonalizable** $\iff \sum_i \alpha_i = \sum_i \zeta_i = n$.

Orthogonally Diagonalizable and Symmetric Matrix

Theorem. $\mathbf{A} \in \mathbb{R}^{n \times n}$ is orthogonally diagonalizable $\iff \mathbf{A}$ is symmetric.

- ▶ **Question.** How to find \mathbf{P} (thus \mathbf{D})?
- ▶ **Spectral Theorem.** If $\mathbf{A} \in \mathbb{R}^{n \times n}$ is symmetric,
 - (a) the eigenvalues are all real
 - (b) the eigenvectors to different eigenvalues are perpendicular.
 - (c) there exists an orthogonal eigenbasis
- ▶ For (c), from each set of eigenvectors, say $\{\mathbf{x}_1, \dots, \mathbf{x}_k\}$ associated with a particular eigenvalue, say λ_j , we can construct another set of eigenvectors $\{\mathbf{x}'_1, \dots, \mathbf{x}'_k\}$ that are orthonormal, using the Gram–Schmidt process.
- ▶ Then, all eigenvectors can form an orthonormal basis.

Example

► **Example.** $\mathbf{A} = \begin{pmatrix} 3 & 2 & 2 \\ 2 & 3 & 2 \\ 2 & 2 & 3 \end{pmatrix}$. $p_{\mathbf{A}}(\lambda) = -(\lambda - 1)^2(\lambda - 7)$, thus $\lambda_1 = 1, \lambda_2 = 7$

$$E_1 = \text{span} \left(\begin{pmatrix} -1 \\ 1 \\ 0 \end{pmatrix}, \begin{pmatrix} -1 \\ 0 \\ 1 \end{pmatrix} \right), \quad E_7 = \text{span} \left(\begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix} \right)$$

- $(111)^{\top}$ is perpendicular to $(-110)^{\top}$ and $(-101)^{\top}$
- $\begin{pmatrix} -1 \\ 1 \\ 0 \end{pmatrix}$ and $\begin{pmatrix} -1/2 \\ -1/2 \\ 1 \end{pmatrix}$ (for $\lambda = 1$) and $\begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}$ (for $\lambda = 7$) are the orthogonal basis in \mathbb{R}^3 .
- After normalization, we can make the orthonormal basis.

Eigendecomposition

- ▶ **Theorem.** The following are equivalent.
 - (a) A square matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ can be factorized into $\mathbf{A} = \mathbf{P}\mathbf{D}\mathbf{P}^{-1}$, where $\mathbf{P} \in \mathbb{R}^{n \times n}$ and \mathbf{D} is the diagonal matrix whose diagonal entries are eigenvalues of \mathbf{A} .
 - (b) The eigenvectors of \mathbf{A} form a basis of \mathbb{R}^n (i.e., the n eigenvectors of \mathbf{A} are linearly independent)
- ▶ The above implies the columns of \mathbf{P} are the n eigenvectors of \mathbf{A} (because $\mathbf{AP} = \mathbf{PD}$)
- ▶ When \mathbf{A} is symmetric, \mathbf{P} is an orthogonal matrix, so $\mathbf{P}^\top = \mathbf{P}^{-1}$
- ▶ If \mathbf{A} is symmetric, then (b) holds (Spectral Theorem).

Example of Orthogonal Diagonalization (1)

- ▶ Eigendecomposition for $\mathbf{A} = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$
- ▶ Eigenvalues: $\lambda_1 = 1, \lambda_2 = 3$
- ▶ (normalized) eigenvectors: $\mathbf{p}_1 = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 \\ -1 \end{pmatrix}, \mathbf{p}_2 = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 \\ 1 \end{pmatrix}$.
- ▶ \mathbf{p}_1 and \mathbf{p}_2 linearly independent, so \mathbf{A} is diagonalizable.
- ▶ $\mathbf{P} = (\mathbf{p}_1 \quad \mathbf{p}_2) = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ -1 & 1 \end{pmatrix}$
- ▶ $\mathbf{D} = \mathbf{P}^{-1} \mathbf{A} \mathbf{P} = \begin{pmatrix} 1 & 0 \\ 0 & 3 \end{pmatrix}$. Finally, we get $\mathbf{A} = \mathbf{P} \mathbf{D} \mathbf{P}^{-1}$

Example of Orthogonal Diagonalization (2)

► $\mathbf{A} = \begin{pmatrix} 1 & 2 & 2 \\ 2 & 1 & 2 \\ 2 & 2 & 1 \end{pmatrix}$

► Eigenvalues: $\lambda_1 = -1$, $\lambda_2 = 5$
(algebraic mult.: 2, 1)

► $E_{-1} = \text{span} \begin{pmatrix} -1 \\ 1 \\ 0 \end{pmatrix}, \begin{pmatrix} -1 \\ 0 \\ 1 \end{pmatrix}$

► Gram–Schmidt ⇒

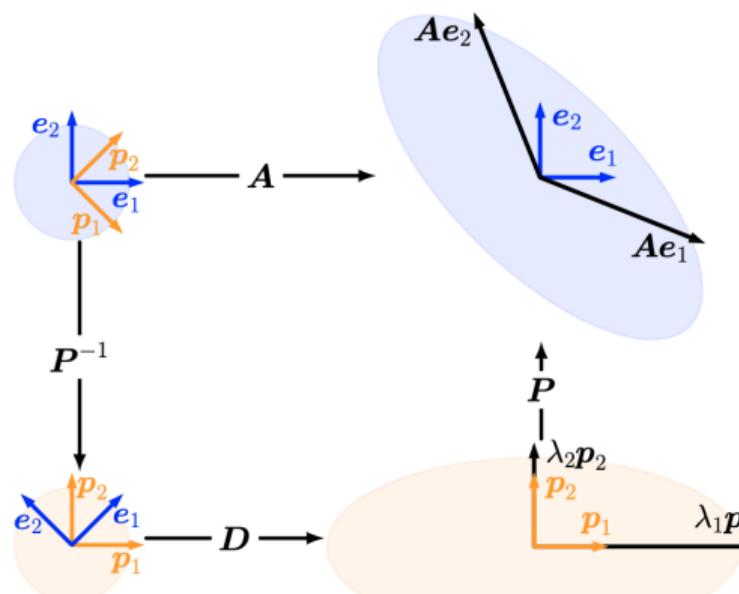
$$\text{span} \frac{1}{\sqrt{2}} \begin{pmatrix} -1 \\ 1 \\ 0 \end{pmatrix}, \frac{1}{\sqrt{6}} \begin{pmatrix} -1 \\ -1 \\ 2 \end{pmatrix}$$

► $E_5 = \text{span} \frac{1}{\sqrt{3}} \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}$

► $\mathbf{P} = \begin{pmatrix} -1/\sqrt{2} & -1/\sqrt{6} & 1/\sqrt{3} \\ 1/\sqrt{2} & -1/\sqrt{6} & 1/\sqrt{3} \\ 0 & 2/\sqrt{6} & 1/\sqrt{3} \end{pmatrix}$

► $\mathbf{D} = \mathbf{P}^\top \mathbf{A} \mathbf{P} = \begin{pmatrix} -1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 5 \end{pmatrix}$

Eigendecomposition: Geometric Interpretation



Question. Can we generalize this beautiful result to a general matrix $A \in \mathbb{R}^{m \times n}$?

Quick Recap: Eigendecomposition

What we just learned:

- ▶ **Eigendecomposition:** $\mathbf{A} = \mathbf{P}\mathbf{D}\mathbf{P}^{-1}$ — change basis, scale, change back
- ▶ Works when \mathbf{A} is square and has enough independent eigenvectors
- ▶ For **symmetric** matrices: \mathbf{P} is orthogonal ($\mathbf{P}^{-1} = \mathbf{P}^{\top}$)
- ▶ **Problem:** What if \mathbf{A} is not square? We need something more general...
- ▶ **Next up:** SVD — works for *any* matrix!

Roadmap

- (1) Determinant and Trace
- (2) Eigenvalues and Eigenvectors
- (3) Cholesky Decomposition
- (4) Eigendecomposition and Diagonalization
- (5) Singular Value Decomposition
- (6) Matrix Approximation
- (7) Matrix Phylogeny

Why Should I Care About SVD?

- ▶ **Image compression:** A 1000×1000 image can be stored using only 50 singular values with almost no visible quality loss
- ▶ **Recommendation systems:** Netflix and Spotify use SVD to predict what you might like
- ▶ **Natural Language Processing:** SVD powers Latent Semantic Analysis — understanding meaning from text
- ▶ **The universal tool:** Unlike eigendecomposition, SVD works for **any** matrix (any shape, any type)

In one sentence: SVD is the Swiss Army knife of linear algebra — it works everywhere.

Storyline

- ▶ Eigendecomposition (also called EVD: EigenValue Decomposition): (Orthogonal) Diagonalization for symmetric matrices $\mathbf{A} \in \mathbb{R}^{n \times n}$.
- ▶ Extensions: Singular Value Decomposition (SVD)
 - First extension: diagonalization for non-symmetric, but still square matrices $\mathbf{A} \in \mathbb{R}^{n \times n}$
 - Second extension: diagonalization for non-symmetric, and non-square matrices $\mathbf{A} \in \mathbb{R}^{m \times n}$
- ▶ **Background.** For $\mathbf{A} \in \mathbb{R}^{m \times n}$, a matrix $\mathbf{S} \stackrel{\text{def}}{=} \mathbf{A}^\top \mathbf{A} \in \mathbb{R}^{n \times n}$ is always symmetric, positive semidefinite.
 - ▶ Symmetric, because $\mathbf{S}^\top = (\mathbf{A}^\top \mathbf{A})^\top = \mathbf{A}^\top \mathbf{A} = \mathbf{S}$.
 - ▶ Positive semidefinite, because $\mathbf{x}^\top \mathbf{S} \mathbf{x} = \mathbf{x}^\top \mathbf{A}^\top \mathbf{A} \mathbf{x} = (\mathbf{Ax})^\top (\mathbf{Ax}) \geq 0$.
 - ▶ If $\text{rank}(\mathbf{A}) = n$, then \mathbf{S} is symmetric and positive definite.

Singular Value Decomposition

- **Theorem.** $\mathbf{A} \in \mathbb{R}^{m \times n}$ with rank $r \in [0, \min(m, n)]$. The **SVD** of \mathbf{A} is a decomposition of the form

$$\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^\top,$$
$$\begin{matrix} & n \\ \approx & \boxed{\mathbf{A}} \end{matrix} = \begin{matrix} & m \\ \approx & \boxed{\mathbf{U}} \end{matrix} \begin{matrix} & n \\ \approx & \boxed{\Sigma} \end{matrix} \begin{matrix} & n \\ \approx & \boxed{\mathbf{V}^\top} \end{matrix} \approx$$

with an orthogonal matrix $\mathbf{U} = (\mathbf{u}_1 \ \cdots \ \mathbf{u}_m) \in \mathbb{R}^{m \times m}$ and an orthogonal matrix $\mathbf{V} = (\mathbf{v}_1 \ \cdots \ \mathbf{v}_n) \in \mathbb{R}^{n \times n}$. Moreover, Σ is an $m \times n$ matrix with $\Sigma_{ii} = \sigma_i \geq 0$ and $\Sigma_{ij} = 0$, $i \neq j$, which is uniquely determined for \mathbf{A} .

- Note
- The diagonal entries σ_i , $i = 1, \dots, r$ are called **singular values**.
 - \mathbf{u}_i and \mathbf{v}_j are called **left** and **right singular vectors**, respectively.

SVD: The Geometric Picture

- ▶ Every matrix transformation can be broken into **three simple steps**:

$$\mathbf{A} = \underbrace{\mathbf{U}}_{\text{Rotate}} \underbrace{\Sigma}_{\text{Stretch}} \underbrace{\mathbf{V}^T}_{\text{Rotate}}$$

- ▶ **Step 1 (\mathbf{V}^T)**: Rotate/align the input to special directions
- ▶ **Step 2 (Σ)**: Stretch each direction by its singular value σ_i
- ▶ **Step 3 (\mathbf{U})**: Rotate the result into the output space
- ▶ **Analogy**: Think of reshaping clay:
 - ▶ First, orient the clay along its natural axes
 - ▶ Then, squeeze or stretch along each axis
 - ▶ Finally, rotate to the desired position

SVD: How It Works (for $\mathbf{A} \in \mathbb{R}^{n \times n}$)

- $\mathbf{A} \in \mathbb{R}^{n \times n}$ with rank $r \leq n$. Then, $\mathbf{A}^\top \mathbf{A}$ is symmetric.
- Orthogonal diagonalization of $\mathbf{A}^\top \mathbf{A}$:

$$\mathbf{A}^\top \mathbf{A} = \mathbf{V} \mathbf{D} \mathbf{V}^\top.$$

- $\mathbf{D} = \begin{pmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_n \end{pmatrix}$ and an orthogonal matrix $\mathbf{V} = (\mathbf{v}_1 \cdots \mathbf{v}_n)$, where $\lambda_1 \geq \cdots \geq \lambda_r > \lambda_{r+1} = \cdots = \lambda_n = 0$ are the eigenvalues of $\mathbf{A}^\top \mathbf{A}$ and $\{\mathbf{v}_i\}$ are orthonormal.
- All λ_i are non-negative: for eigenvector \mathbf{v}_i with eigenvalue λ_i ,

$$0 \leq \|\mathbf{A}\mathbf{v}_i\|^2 = \mathbf{v}_i^\top \mathbf{A}^\top \mathbf{A} \mathbf{v}_i = \lambda_i \|\mathbf{v}_i\|^2 \implies \lambda_i \geq 0$$

- $\text{rank}(\mathbf{A}) = \text{rank}(\mathbf{A}^\top \mathbf{A}) = \text{rank}(\mathbf{D}) = r$

- Choose $\mathbf{U}' = (\mathbf{u}_1 \cdots \mathbf{u}_r)$, where

$$\mathbf{u}_i = \frac{\mathbf{A}\mathbf{v}_i}{\sqrt{\lambda_i}}, \quad 1 \leq i \leq r.$$

- We can construct $\{\mathbf{u}_i\}$, $i = r+1, \dots, n$, so that $\mathbf{U} = (\mathbf{u}_1 \cdots \mathbf{u}_n)$ is an orthonormal basis of \mathbb{R}^n .

Example

► $\mathbf{A} = \begin{pmatrix} 1 & 0 & 1 \\ -2 & 1 & 0 \end{pmatrix}$

► $\mathbf{A}^\top \mathbf{A} = \begin{pmatrix} 5 & -2 & 1 \\ -2 & 1 & 0 \\ 1 & 0 & 1 \end{pmatrix} = \mathbf{V} \mathbf{D} \mathbf{V}^\top,$

$$\mathbf{D} = \begin{pmatrix} 6 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{pmatrix}, \mathbf{V} = \begin{pmatrix} \frac{5}{\sqrt{30}} & \frac{-2}{\sqrt{30}} & \frac{1}{\sqrt{30}} \\ 0 & \frac{1}{\sqrt{5}} & \frac{2}{\sqrt{5}} \\ \frac{-1}{\sqrt{6}} & \frac{-2}{\sqrt{6}} & \frac{1}{\sqrt{6}} \end{pmatrix}$$

► $\text{rank}(\mathbf{A}) = 2$ because we have two singular values $\sigma_1 = \sqrt{6}$ and $\sigma_2 = 1$

► $\Sigma = \begin{pmatrix} \sqrt{6} & 0 & 0 \\ 0 & 1 & 0 \end{pmatrix}$

► $\mathbf{u}_1 = \mathbf{A}\mathbf{v}_1/\sigma_1 = \begin{pmatrix} \frac{1}{\sqrt{5}} \\ \frac{-2}{\sqrt{5}} \end{pmatrix}$

► $\mathbf{u}_2 = \mathbf{A}\mathbf{v}_2/\sigma_2 = \begin{pmatrix} \frac{2}{\sqrt{5}} \\ \frac{1}{\sqrt{5}} \end{pmatrix}$

► $\mathbf{U} = (\mathbf{u}_1 \quad \mathbf{u}_2) = \frac{1}{\sqrt{5}} \begin{pmatrix} 1 & 2 \\ -2 & 1 \end{pmatrix}$

► Then, we can see that $\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^\top$.

EVD ($\mathbf{A} = \mathbf{P}\mathbf{D}\mathbf{P}^{-1}$) vs. SVD ($\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^\top$)

- ▶ SVD: always exists, EVD: square matrix and exists if we can find a basis of eigenvectors (such as symmetric matrices)
- ▶ \mathbf{P} in EVD is not necessarily orthogonal (only true for symmetric \mathbf{A}), but \mathbf{U} and \mathbf{V} are orthogonal (so representing rotations)
- ▶ Both EVD and SVD: (i) basis change in the domain, (ii) independent scaling of each new basis vector and mapping from domain to codomain, (iii) basis change in the codomain. The difference: for SVD, different vector spaces of domain and codomain.
- ▶ SVD and EVD are closely related through their projections
 - ▶ The left-singular (resp. right-singular) vectors of \mathbf{A} are eigenvectors of $\mathbf{A}\mathbf{A}^\top$ (resp. $\mathbf{A}^\top\mathbf{A}$)
 - ▶ The singular values of \mathbf{A} are the square roots of eigenvalues of $\mathbf{A}\mathbf{A}^\top$ and $\mathbf{A}^\top\mathbf{A}$
 - ▶ When \mathbf{A} is symmetric, EVD = SVD (from spectral theorem)

Different Forms of SVD

- When $\text{rank}(\mathbf{A}) = r$, we can construct SVD as the following with only non-zero diagonal entries in Σ :

$$\mathbf{A} = \overbrace{\mathbf{U}}^{m \times r} \overbrace{\Sigma}^{r \times r} \overbrace{\mathbf{V}^T}^{r \times n}$$

- We can even truncate the decomposed matrices, which can be an approximation of \mathbf{A} : for $k < r$

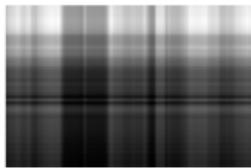
$$\mathbf{A} \approx \overbrace{\mathbf{U}}^{m \times k} \overbrace{\Sigma}^{k \times k} \overbrace{\mathbf{V}^T}^{k \times n}$$

We will cover this in the next slides.

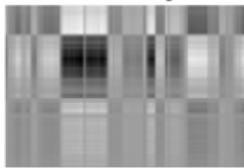
Matrix Approximation via SVD



(a) Original image \mathbf{A} .



(b) \mathbf{A}_1 , $\sigma_1 \approx 228,052$.



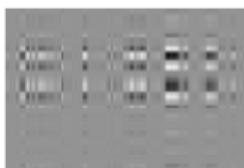
(c) \mathbf{A}_2 , $\sigma_2 \approx 40,647$.



(d) \mathbf{A}_3 , $\sigma_3 \approx 26,125$.



(e) \mathbf{A}_4 , $\sigma_4 \approx 20,232$.



(f) \mathbf{A}_5 , $\sigma_5 \approx 15,436$.

- ▶ $\mathbf{A} = \sum_{i=1}^r \sigma_i \underbrace{\mathbf{u}_i \mathbf{v}_i^\top}_{\mathbf{A}_i}$, where \mathbf{A}_i is the outer product³ of \mathbf{u}_i and \mathbf{v}_i
- ▶ Rank k -approximation: $\hat{\mathbf{A}}(k) = \sum_{i=1}^k \sigma_i \mathbf{A}_i$, $k < r$

³If \mathbf{u} and \mathbf{v} are both nonzero, then the outer product matrix $\mathbf{u}\mathbf{v}^\top$ always has matrix rank 1. Indeed, the columns of the outer product are all proportional to the first column.

How Close $\hat{\mathbf{A}}(k)$ is to \mathbf{A} ?

- ▶ **Definition.** Spectral Norm of a Matrix. For $\mathbf{A} \in \mathbb{R}^{m \times n}$, $\|\mathbf{A}\|_2 \stackrel{\text{def}}{=} \max_{\mathbf{x}} \frac{\|\mathbf{Ax}\|_2}{\|\mathbf{x}\|_2}$
 - ▶ As a concept of length of \mathbf{A} , it measures how long any vector \mathbf{x} can at most become, when multiplied by \mathbf{A}
- ▶ **Theorem.** Eckart–Young. For $\mathbf{A} \in \mathbb{R}^{m \times n}$ of rank r and $\mathbf{B} \in \mathbb{R}^{m \times n}$ of rank k , for any $k \leq r$, we have:

$$\hat{\mathbf{A}}(k) = \arg \min_{\text{rank}(\mathbf{B})=k} \|\mathbf{A} - \mathbf{B}\|_2, \quad \text{and} \quad \|\mathbf{A} - \hat{\mathbf{A}}(k)\|_2 = \sigma_{k+1}$$

- ▶ Quantifies how much error is introduced by the SVD-based approximation
- ▶ $\hat{\mathbf{A}}(k)$ is optimal in the sense that such SVD-based approximation is the best one among all rank- k approximations.
- ▶ In other words, it is a projection of the full-rank matrix \mathbf{A} onto a lower-dimensional space of rank-at-most- k matrices.

Quick Recap: SVD & Matrix Approximation

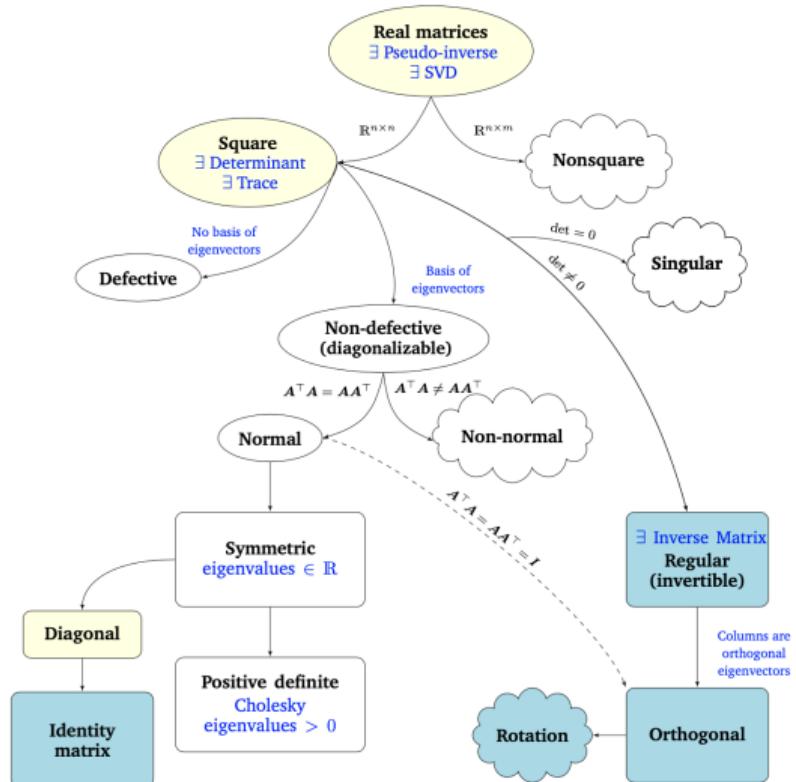
What we just learned:

- ▶ **SVD:** $\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^\top$ — rotate, stretch, rotate (works for *any* matrix)
- ▶ Singular values tell you how important each direction is
- ▶ **Truncated SVD** gives the best rank- k approximation (Eckart–Young theorem)
- ▶ **Applications:** image compression, dimensionality reduction, noise removal

Roadmap

- (1) Determinant and Trace
- (2) Eigenvalues and Eigenvectors
- (3) Cholesky Decomposition
- (4) Eigendecomposition and Diagonalization
- (5) Singular Value Decomposition
- (6) Matrix Approximation
- (7) **Matrix Phylogeny**

Phylogenetic Tree of Matrices



Which Decomposition Should I Use?

	Requirement	Form	Best For
LU	Any square matrix	$\mathbf{A} = \mathbf{L}\mathbf{U}$	Solving $\mathbf{Ax} = \mathbf{b}$
Cholesky	Symmetric + PD	$\mathbf{A} = \mathbf{LL}^\top$	Covariance, fast solving
Eigendecompr.	Square + basis of eigvecs	$\mathbf{A} = \mathbf{PDP}^{-1}$	Powers \mathbf{A}^k , stability
SVD	Any matrix!	$\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^\top$	Approximation, compression

Rule of thumb: Start with SVD — it always works.

Use Cholesky or Eigendecomposition when you know your matrix is special (symmetric, PD, etc.)

Common Mistakes to Avoid

(1) Applying Eigendecomposition to non-square matrices

Eigendecomposition requires $\mathbf{A} \in \mathbb{R}^{n \times n}$. For rectangular matrices, use SVD.

(2) Confusing eigenvalues with singular values

Eigenvalues can be negative or complex. Singular values are always ≥ 0 .

(3) Forgetting to check symmetry before Cholesky

Cholesky requires \mathbf{A} to be symmetric *and* positive definite.

(4) Assuming eigenvectors are unique

Any scalar multiple of an eigenvector is also an eigenvector.

(5) Mixing up algebraic and geometric multiplicity

Algebraic = root count in $p_{\mathbf{A}}(\lambda)$. Geometric = dimension of eigenspace.

Key Takeaways

- (1) Determinant & Trace summarize a matrix: $\det(\mathbf{A}) = \prod_i \lambda_i$, $\text{tr}(\mathbf{A}) = \sum_i \lambda_i$
- (2) Eigendecomposition: $\mathbf{A} = \mathbf{P}\mathbf{D}\mathbf{P}^{-1}$ reveals how a matrix stretches space along its eigenvectors
- (3) Cholesky: $\mathbf{A} = \mathbf{L}\mathbf{L}^\top$ for SPD matrices — fast, stable factorization
- (4) SVD: $\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^\top$ works for *any* matrix — the most general decomposition
- (5) Matrix Approximation: Truncated SVD gives the optimal rank- k approximation (Eckart–Young)
- (6) Concept Chain:

Eigenvalues → Diagonalization → SVD → Low-Rank Approximation

Questions?

Review Questions (1): Determinant Basics

- 1) Compute $\det(\mathbf{A})$ for $\mathbf{A} = \begin{pmatrix} 2 & -1 \\ 4 & 3 \end{pmatrix}$ and decide if \mathbf{A} is invertible.
- 2) Let \mathbf{A} be upper triangular with diagonal entries $(3, -2, 5, 1)$. Compute $\det(\mathbf{A})$.
- 3) True/False (justify): If $\det(\mathbf{A}) = 0$ then the columns of \mathbf{A} are linearly dependent.
- 4) Use Laplace expansion along the *first row* to write $\det(\mathbf{A})$ for a 3×3 matrix symbolically.

Review Questions (2): Determinant Properties + Elimination

- 1) Suppose \mathbf{B} is obtained from \mathbf{A} by swapping two rows. How does $\det(\mathbf{B})$ relate to $\det(\mathbf{A})$?
- 2) Suppose \mathbf{B} is obtained from \mathbf{A} by multiplying one row by α . How does $\det(\mathbf{B})$ change?
- 3) Given $\det(\mathbf{A}) = 5$ and $\det(\mathbf{B}) = -2$, compute:
 - ▶ $\det(\mathbf{AB})$
 - ▶ $\det(\mathbf{A}^T)$
 - ▶ $\det(\mathbf{A}^{-1})$
- 4) Explain (in one sentence) why Gaussian elimination can be used to compute determinants efficiently.

Review Questions (3): Trace + Cyclic Invariance

- 1) Compute $\text{tr}(\mathbf{A})$ for $\mathbf{A} = \begin{pmatrix} 1 & 2 & 3 \\ 0 & -1 & 4 \\ 5 & 0 & 2 \end{pmatrix}$.
- 2) True/False (justify): $\text{tr}(\mathbf{AB}) = \text{tr}(\mathbf{A})\text{tr}(\mathbf{B})$.
- 3) Prove using cyclic property:

$$\text{tr}(\mathbf{AB}) = \text{tr}(\mathbf{BA}) \quad (\mathbf{A} \in \mathbb{R}^{n \times k}, \mathbf{B} \in \mathbb{R}^{k \times n}).$$

- 4) Let $\mathbf{B} = \mathbf{S}^{-1}\mathbf{AS}$ (similar matrices). Show $\text{tr}(\mathbf{B}) = \text{tr}(\mathbf{A})$.

Review Questions (4): Characteristic Polynomial & Eigenvalues

- 1) For $\mathbf{A} = \begin{pmatrix} 4 & 2 \\ 1 & 3 \end{pmatrix}$ compute $p_{\mathbf{A}}(\lambda) = \det(\mathbf{A} - \lambda\mathbf{I})$.
- 2) Solve $p_{\mathbf{A}}(\lambda) = 0$ to obtain eigenvalues.
- 3) For each eigenvalue, find a basis of its eigenspace E_{λ} .
- 4) Explain the equivalence:

$$\det(\mathbf{A} - \lambda\mathbf{I}) = 0 \iff (\mathbf{A} - \lambda\mathbf{I})\mathbf{x} = \mathbf{0} \text{ has a nontrivial solution.}$$

Review Questions (5): Eigen Geometry + Invariants

- 1) Give a geometric meaning of an eigenvector/eigenvalue pair (λ, \mathbf{x}) .
- 2) For $\mathbf{A} = \begin{pmatrix} 1 & \frac{1}{2} \\ 0 & 1 \end{pmatrix}$:
 - ▶ Find eigenvalues and eigenvectors.
 - ▶ Describe the mapping (rotation/shear/stretch/collapse).
- 3) If $\mathbf{A}' = \mathbf{S}^{-1}\mathbf{AS}$, show \mathbf{A} and \mathbf{A}' have the same eigenvalues.
- 4) Using eigenvalues $\{\lambda_i\}_{i=1}^n$, state formulas for $\det(\mathbf{A})$ and $\text{tr}(\mathbf{A})$.

Review Questions (6): LU & Cholesky

- 1) Conceptual: In Gaussian elimination, why do elimination matrices \mathbf{E}_i tend to be lower triangular?
- 2) If $\mathbf{A} = \mathbf{LU}$ (LU), what are the shapes/properties of \mathbf{L} and \mathbf{U} ?
- 3) State the condition under which Cholesky exists, and write the factorization form.
- 4) Let \mathbf{A} be SPD and $\mathbf{A} = \mathbf{LL}^\top$. Show:

$$\det(\mathbf{A}) = \det(\mathbf{L})^2 = \prod_i \ell_{ii}^2.$$

Review Questions (7): Diagonalization & Multiplicities

- 1) Define: diagonalizable matrix. What does it mean that \mathbf{A} is similar to a diagonal matrix?
- 2) For $\mathbf{A} = \begin{pmatrix} 2 & 1 \\ 0 & 2 \end{pmatrix}$:
 - ▶ Find eigenvalues and eigenvectors.
 - ▶ Compute algebraic and geometric multiplicities.
 - ▶ Decide if \mathbf{A} is diagonalizable.
- 3) State the diagonalizability criterion in terms of $\sum_i \alpha_i$ and $\sum_i \zeta_i$.
- 4) If $\mathbf{A} = \mathbf{P}\mathbf{D}\mathbf{P}^{-1}$, show $\mathbf{A}^k = \mathbf{P}\mathbf{D}^k\mathbf{P}^{-1}$.

Review Questions (8): Spectral Theorem & Orthogonal Diagonalization

- 1) State the Spectral Theorem for real symmetric matrices.
- 2) Prove/justify: \mathbf{A} is orthogonally diagonalizable $\iff \mathbf{A}$ is symmetric.
- 3) Given eigenvectors for a repeated eigenvalue, explain how Gram–Schmidt is used to build an orthonormal eigenbasis.
- 4) For a symmetric \mathbf{A} , explain why $\mathbf{P}^{-1} = \mathbf{P}^T$ in $\mathbf{A} = \mathbf{P}\mathbf{D}\mathbf{P}^T$.

Review Questions (9): SVD Core Relationships

- 1) State the SVD theorem: $\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^\top$ and define singular values/vectors.
- 2) Show that $\mathbf{A}^\top\mathbf{A}$ is symmetric PSD and explain why its eigenvalues are ≥ 0 .
- 3) Show the relationship:
 - ▶ right singular vectors \mathbf{v}_i are eigenvectors of $\mathbf{A}^\top\mathbf{A}$
 - ▶ singular values satisfy $\sigma_i = \sqrt{\lambda_i(\mathbf{A}^\top\mathbf{A})}$
- 4) When does EVD coincide with SVD? (state the condition and reason)

Review Questions (10): Low-Rank Approximation (Eckart–Young)

- 1) Write the rank-1 expansion:

$$\mathbf{A} = \sum_{i=1}^r \sigma_i \mathbf{u}_i \mathbf{v}_i^\top.$$

Explain why each term has rank 1.

- 2) Define the spectral norm:

$$\|\mathbf{A}\|_2 = \max_{\mathbf{x} \neq 0} \frac{\|\mathbf{Ax}\|_2}{\|\mathbf{x}\|_2}.$$

- 3) State Eckart–Young and interpret:

$$\left\| \mathbf{A} - \hat{\mathbf{A}}(k) \right\|_2 = \sigma_{k+1}.$$

- 4) Practical: If singular values are $(10, 3, 1, 0.2, 0)$, what is the best rank-2 approximation error in spectral norm?

Theory Link (1): Data as an Operator in Hilbert Space

- ▶ Let data matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$.
- ▶ Covariance operator:

$$\mathbf{C} = \frac{1}{n} \mathbf{X}^\top \mathbf{X}$$

- ▶ Spectral theorem:

$$\mathbf{C} = \mathbf{V} \Lambda \mathbf{V}^\top$$

- ▶ Eigenvectors = principal directions Eigenvalues = energy along directions.
- ▶ Functional view:

$$\mathbf{Cv}_i = \lambda_i \mathbf{v}_i$$

- ▶ Interpretation:
 - ▶ Data lies near a low-dimensional subspace.
 - ▶ Learning = discovering dominant eigenspaces.
- ▶ Deep networks implicitly approximate:

Projection onto dominant spectral components

Theory Link (2): Low-Rank Hypothesis in Deep Models

- ▶ Layer weight:

$$\mathbf{W} \in \mathbb{R}^{m \times n}$$

- ▶ Empirical phenomenon:

$$\text{rank}(\mathbf{W}) \ll \min(m, n)$$

- ▶ Optimal rank- k approximation:

$$\mathbf{W}_k = \arg \min_{\text{rank}(\mathbf{B})=k} \|\mathbf{W} - \mathbf{B}\|_F$$

- ▶ Eckart–Young:

$$\mathbf{W}_k = \mathbf{U}_k \Sigma_k \mathbf{V}_k^\top$$

- ▶ Meaning:

- ▶ Networks learn compressed linear operators.
- ▶ Redundant directions collapse spectrally.

- ▶ In CNNs:

Filters \Rightarrow learned basis functions

Theory Link (3): Spectral Geometry of Representations

- ▶ Deep features:

$$\mathbf{H} = f(\mathbf{X})$$

- ▶ Feature covariance:

$$\mathbf{C}_H = \frac{1}{n} \mathbf{H}^\top \mathbf{H}$$

- ▶ Spectrum:

$$\lambda_1 \geq \lambda_2 \geq \cdots \geq \lambda_d$$

- ▶ If:

$$\lambda_1 \gg \lambda_2 \gg \cdots$$

Then:

$$\text{rank}(\mathbf{H}) \approx k$$

- ▶ Trace identity:

$$\text{tr}(\mathbf{C}_H) = \sum_i \lambda_i$$

- ▶ Interpretation:



Total variance in learned representation.

Theory Link (4): Spectral View of Optimization Dynamics

- ▶ Local training behavior governed by Hessian:

$$\mathbf{H} = \nabla^2 L(\theta)$$

- ▶ Quadratic approximation:

$$L(\theta + \delta) \approx L(\theta) + \frac{1}{2} \delta^\top \mathbf{H} \delta$$

- ▶ Eigen-decomposition:

$$\mathbf{H} = \mathbf{Q} \Lambda \mathbf{Q}^\top$$

- ▶ Dynamics:

- ▶ Large $\lambda_i \Rightarrow$ sharp curvature
- ▶ Small $\lambda_i \Rightarrow$ flat directions

- ▶ Generalization hypothesis:

Flat minima \Leftrightarrow smaller dominant eigenvalues

- ▶ Training becomes:

Spectral shaping of the loss landscape

Theory Link (5): Operator View of Deep Networks

- ▶ Linear layer:

$$\mathbf{x} \mapsto \mathbf{Wx}$$

- ▶ SVD:

$$\mathbf{W} = \mathbf{U}\Sigma\mathbf{V}^\top$$

- ▶ Interpretation:

- ▶ \mathbf{V}^\top : change of basis (input space)
- ▶ Σ : anisotropic scaling
- ▶ \mathbf{U} : rotation (output space)

- ▶ Thus each layer:

Rotate → Scale → Rotate

- ▶ Deep network:

$$\mathbf{W}_L \cdots \mathbf{W}_2 \mathbf{W}_1$$

- ▶ Theoretical insight:

- ▶ Learning = hierarchical spectral factorization.
- ▶ Vision models learn multi-scale spectral operators.

This section connects matrix decompositions to modern theoretical foundations of deep learning:

- ▶ Neural Tangent Kernel (NTK) spectrum
- ▶ Fisher Information geometry
- ▶ Spectral bias theory
- ▶ Implicit regularization
- ▶ Operator compression view

Mathematical maturity assumed

PhD Theory (1): Neural Tangent Kernel Spectrum

- ▶ Neural network:

$$f_\theta(x)$$

- ▶ Linearization near initialization:

$$f_\theta(x) \approx f_{\theta_0}(x) + \nabla_{\theta} f_{\theta_0}(x)^\top (\theta - \theta_0)$$

- ▶ NTK matrix:

$$K(x_i, x_j) = \nabla_{\theta} f(x_i)^\top \nabla_{\theta} f(x_j)$$

- ▶ Spectral decomposition:

$$K = \mathbf{Q} \Lambda \mathbf{Q}^\top$$

- ▶ Training dynamics:

$$f_t = (I - e^{-\eta \Lambda t}) f^*$$

- ▶ Insight:

- ▶ Large eigenvalues learn faster
- ▶ Small eigenvalues learn slower

- ▶ Deep learning is spectrally biased

PhD Theory (2): Fisher Information Geometry

- ▶ Fisher matrix:

$$\mathbf{F} = \mathbb{E}[\nabla_{\theta} \log p(x|\theta) \nabla_{\theta} \log p(x|\theta)^{\top}]$$

- ▶ Spectral form:

$$\mathbf{F} = \mathbf{U} \Lambda \mathbf{U}^{\top}$$

- ▶ Interpretation:

- ▶ Eigenvectors = sensitive parameter directions
- ▶ Eigenvalues = curvature strength

- ▶ Natural gradient:

$$\theta_{t+1} = \theta_t - \eta \mathbf{F}^{-1} \nabla L$$

- ▶ Theoretical meaning:

Learning occurs along dominant Fisher modes

- ▶ Connection:

- ▶ Compression
- ▶ Pruning
- ▶ Information bottleneck

PhD Theory (3): Spectral Bias of Deep Networks

- ▶ Target function decomposition:

$$f(x) = \sum_i \alpha_i \phi_i(x)$$

- ▶ Where ϕ_i are eigenfunctions of the NTK.

- ▶ Gradient descent solution:

$$\alpha_i(t) = \alpha_i(0) + (1 - e^{-\eta \lambda_i t}) \alpha_i^*$$

- ▶ Consequence:

- ▶ Low-frequency modes learned first
- ▶ High-frequency modes learned later

- ▶ Known as:

Spectral Bias (Frequency Principle)

- ▶ Vision implication:

- ▶ Networks capture global structure first
- ▶ Fine details later

PhD Theory (4): Implicit Low-Rank Regularization

- ▶ Consider linear network:

$$\mathbf{W} = \mathbf{W}_L \cdots \mathbf{W}_1$$

- ▶ Gradient descent solution minimizes:

$$\|\mathbf{W}\|_* \quad (\text{nuclear norm})$$

- ▶ Nuclear norm:

$$\|\mathbf{W}\|_* = \sum_i \sigma_i$$

- ▶ Equivalent to:

Promoting low-rank structure

- ▶ Hence:

- ▶ Deep networks favor compressed operators
- ▶ Rank collapse emerges naturally

- ▶ Explains:

- ▶ Generalization
- ▶ Compression without explicit constraints

PhD Theory (5): Deep Networks as Hierarchical Operators

- ▶ Each layer:

$$\mathbf{W} = \mathbf{U}\Sigma\mathbf{V}^\top$$

- ▶ Full network:

$$\mathbf{W}_L \cdots \mathbf{W}_1 = (\mathbf{U}_L \Sigma_L \mathbf{V}_L^\top) \cdots (\mathbf{U}_1 \Sigma_1 \mathbf{V}_1^\top)$$

- ▶ Interpretation:

- ▶ Successive rotations
- ▶ Anisotropic scalings
- ▶ Subspace projections

- ▶ Depth creates:

Hierarchical spectral filtering

- ▶ Theoretical view:

- ▶ CNNs \approx multi-scale operator decompositions
- ▶ Vision models \approx spectral manifold learners